



# Kubernetes for you

ENSIAS – IT-Holic

2<sup>nd</sup> of April



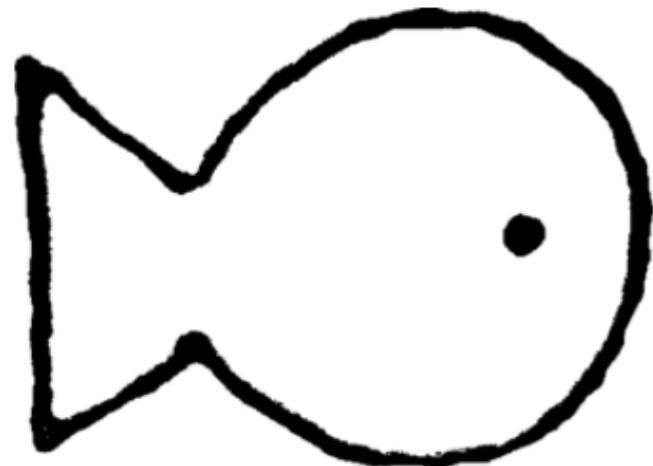
*Karim Benzidane, Ph.D*  
*@The\_Mirak*



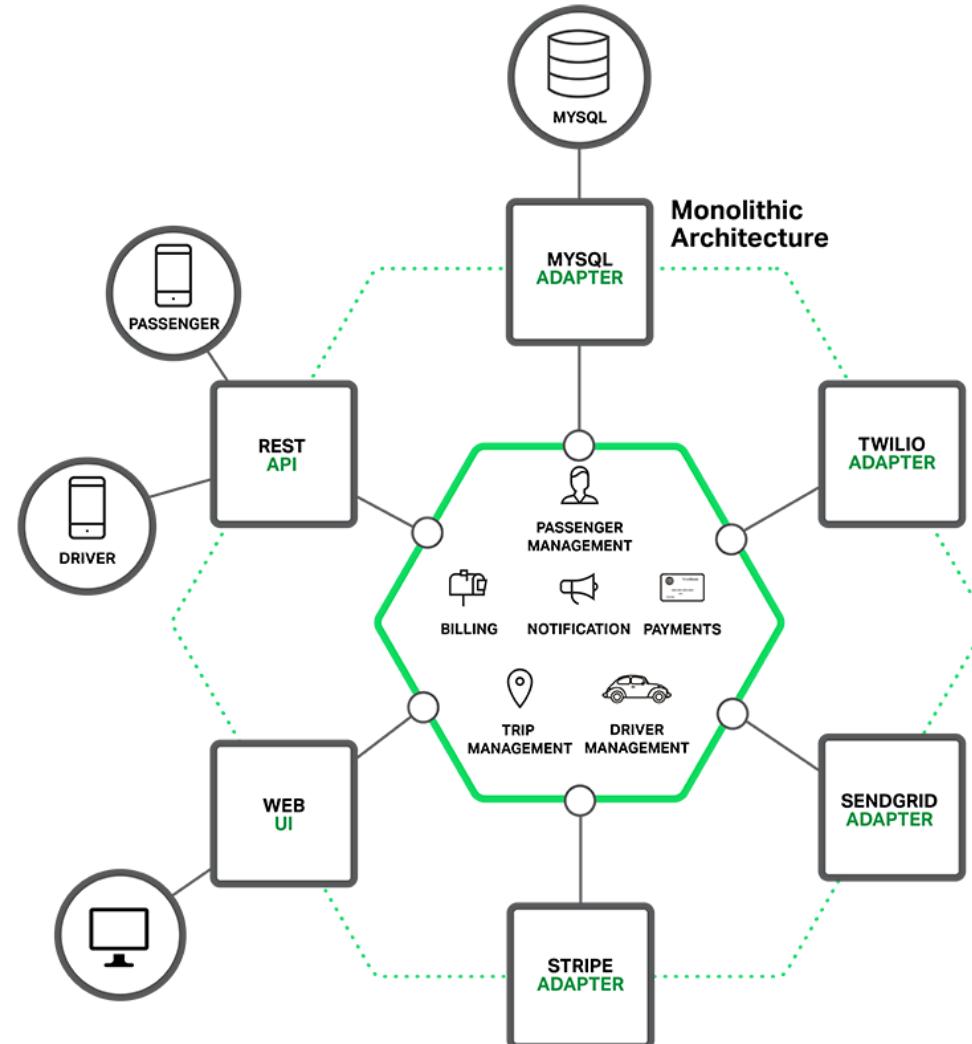
*@cncg\_casa*

Join at  
**slido.com**  
**#316 667**

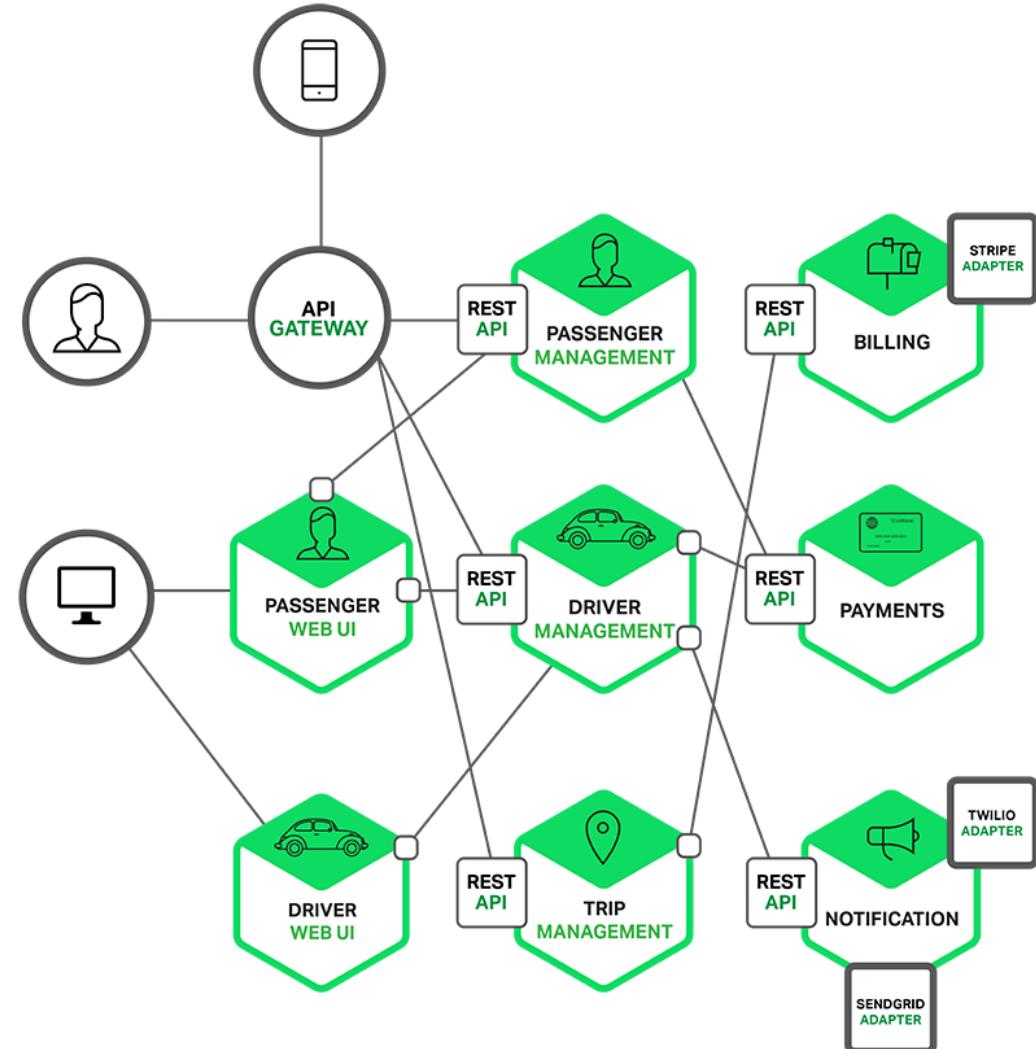
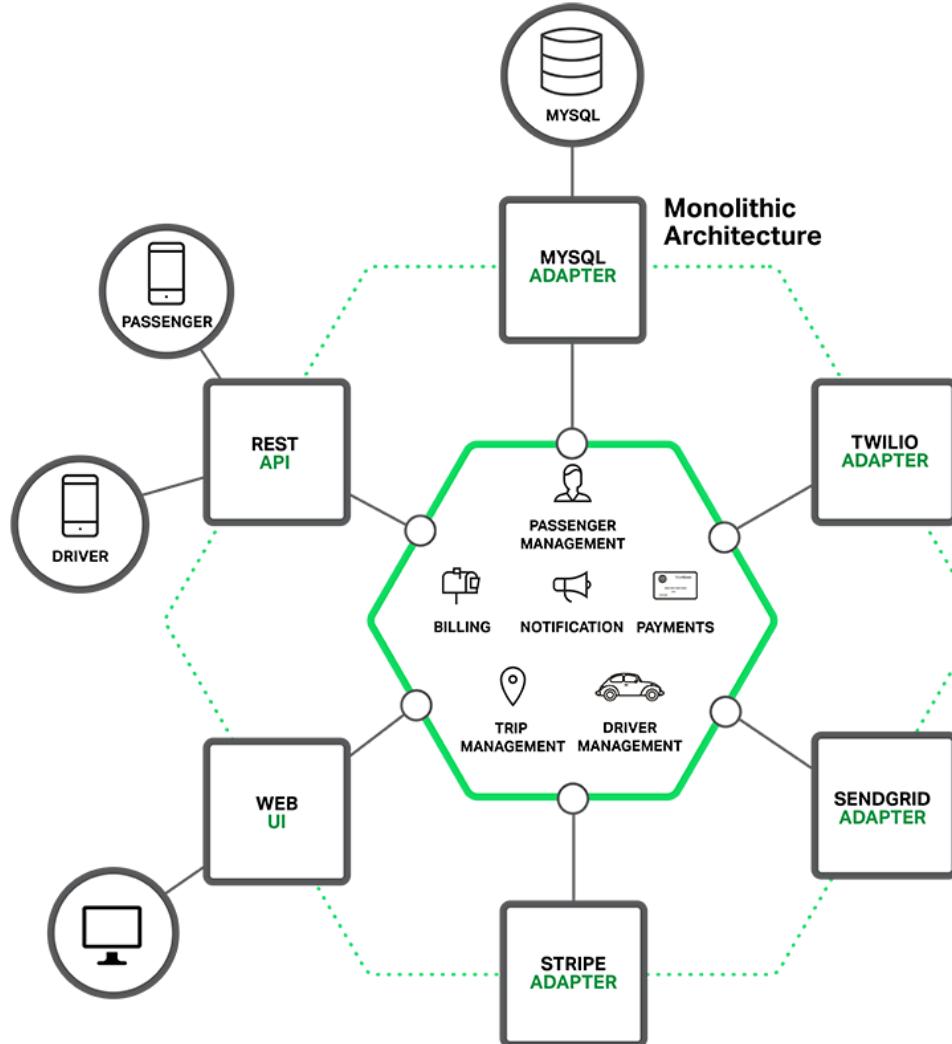




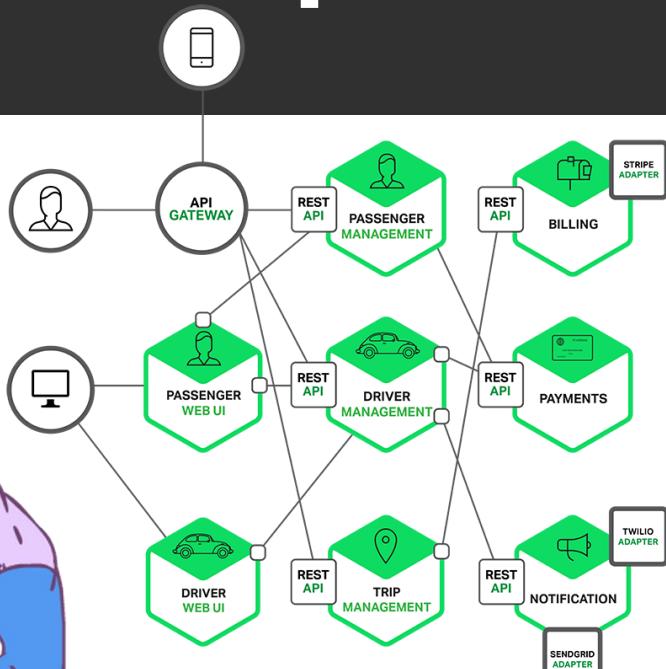
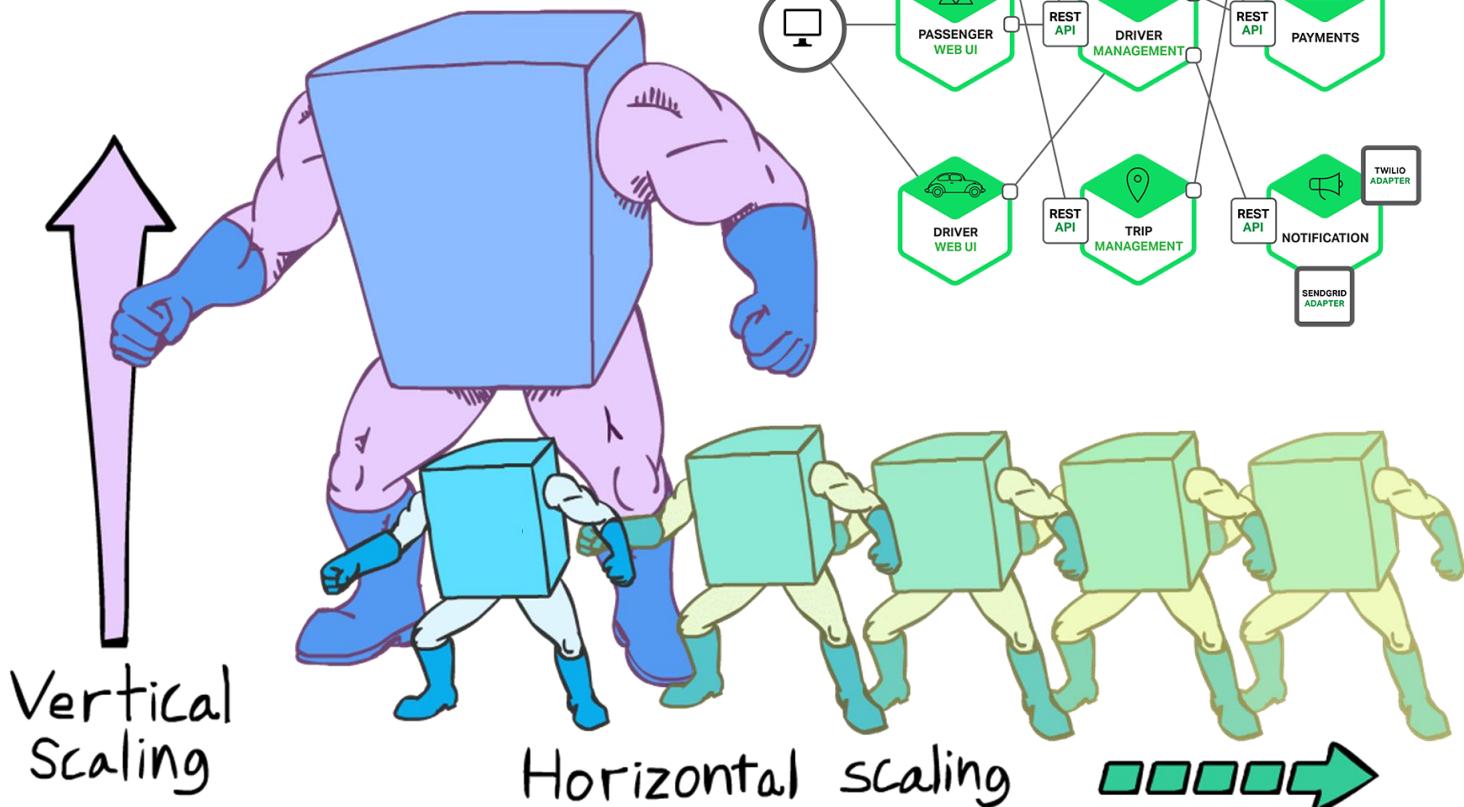
# General Assumption



# General Assumption



# General Assumption



Rachel Dines  
@RachelDines

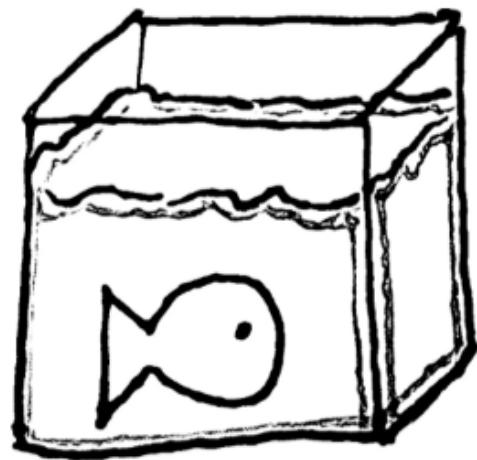
Zoom was adding 5,000-6,000 servers (on AWS) a night as they scaled in the early days of the pandemic 😱 #reInvent

5:47 PM · Dec 15, 2020 · Twitter Web App

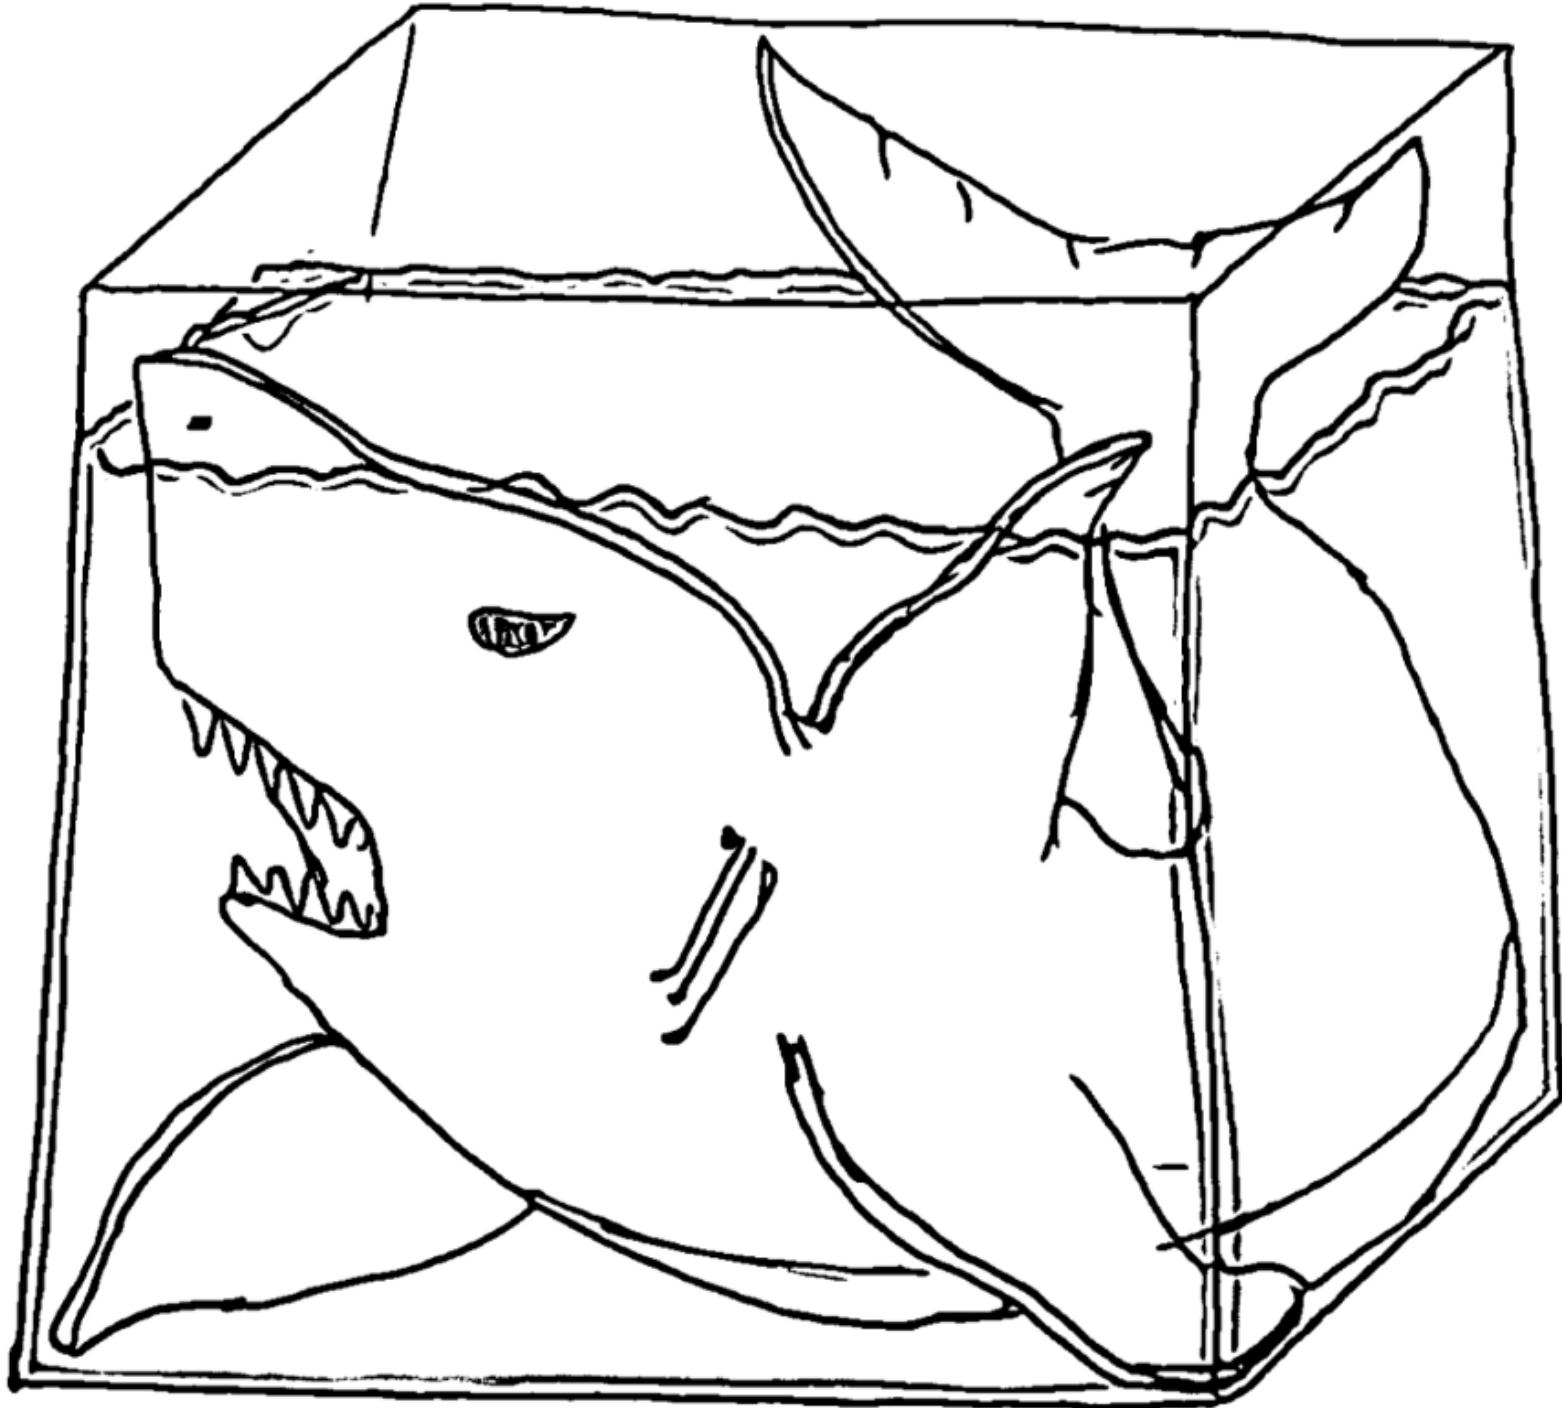
438 Retweets 136 Quote Tweets 3,050 Likes

Tweet your reply

Rachel Dines @RachelDines · Dec 15, 2020  
Replies to @RachelDines  
Zoom was able to adapt to unexpected changes because they broke their app into separate components to do meeting management and video streaming, which have different scaling and latency requirements.  
#reInvent



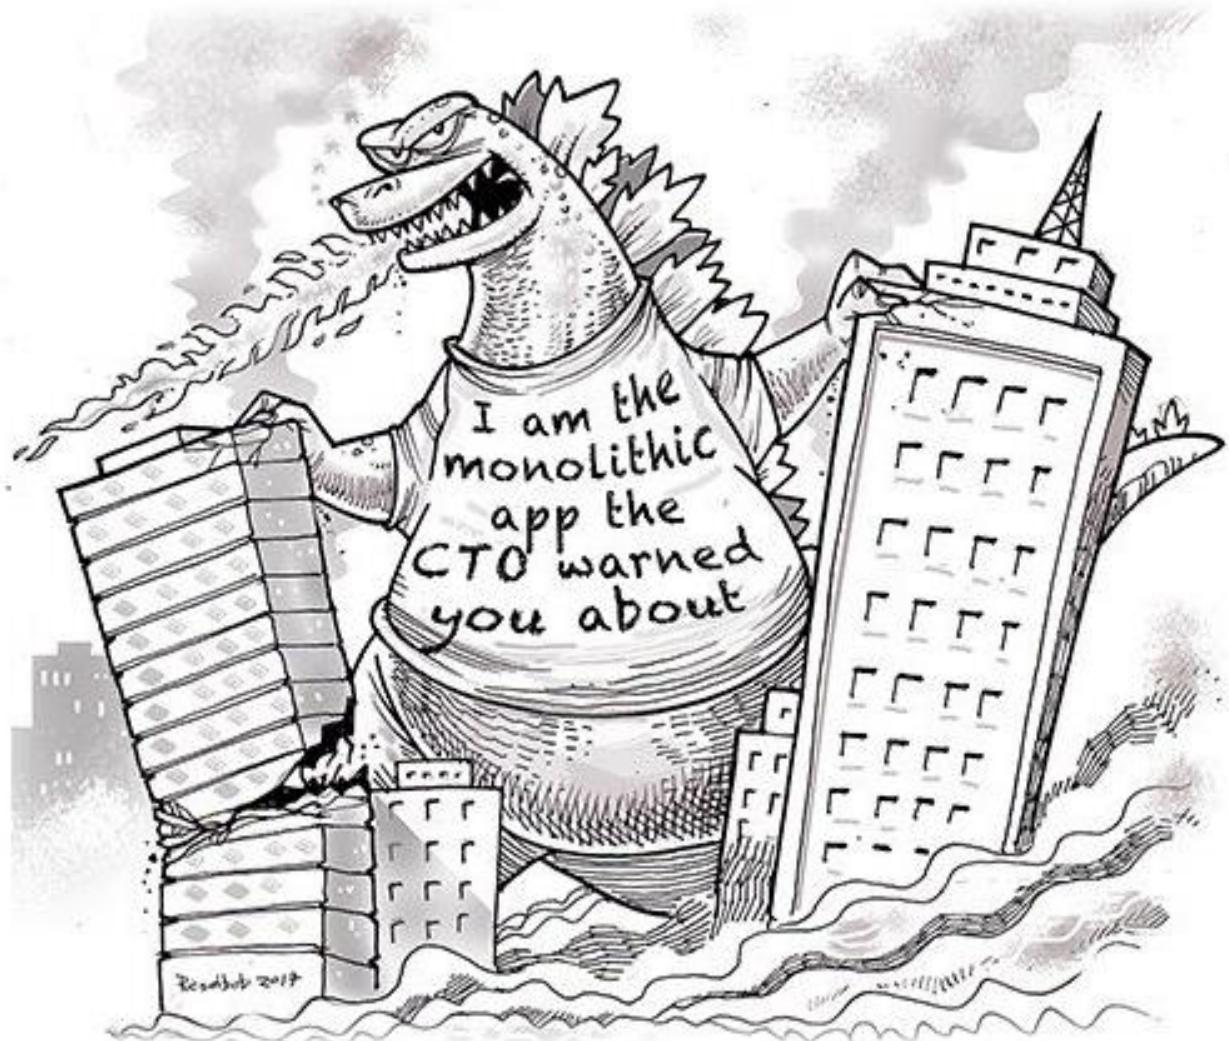
containers



# General Assumption



- As always, microservices also have drawbacks.
- When the **number** of those components **increases**, deployment-related **decisions** become increasingly **difficult**
  - the number of deployment **combinations** increase,
  - the number of **inter-dependencies** **between** the components **increases** by an even greater factor.



# General Assumption



- Microservices need to **find** and **talk to each other**.
- Someone or something **needs to configure** all of them **properly** to **enable** them to **work together** as a single system.
- It becomes increasingly **difficult to configure, manage, and keep the whole system running smoothly**.

# General Assumption



- Ideally, we want **developers** to **deploy** applications **themselves** without having to wait for the **ops** people.
- But **deploying** an application often **requires** an **understanding** of the **underlying infrastructure**.
- Developers **don't** always know those **details**.

# General Assumption



- It's much **harder** to **figure out where to put** each of those **components** to achieve **high resource utilization** and thereby keep the **hardware costs down**.

**Doing all this manually is hard work, Automation is a must**

- This **automation** includes **automatic scheduling** of those **components** to our servers, **automatic configuration**, **supervision**, and **failure-handling**.

In 2006 Google developed  
**cGroups**.

Today, everything at Google  
runs in a container.

On average we launch

**5 billion**

containers every week





In the mid 2000s Google developed **Borg**.



In the mid 2000s Google developed **Borg**.

In 2014 Google developed and open sourced **Kubernetes**



# General Assumption



- After having kept **Borg** and **Omega secret** for a whole decade,
- In **2014** Google introduced **Kubernetes**, an **open-source** system based on the **experience** gained through **Borg**, **Omega**, and other **internal Google systems**.

# General Assumption



*This is where Kubernetes comes in.*



# What is Kubernetes !?

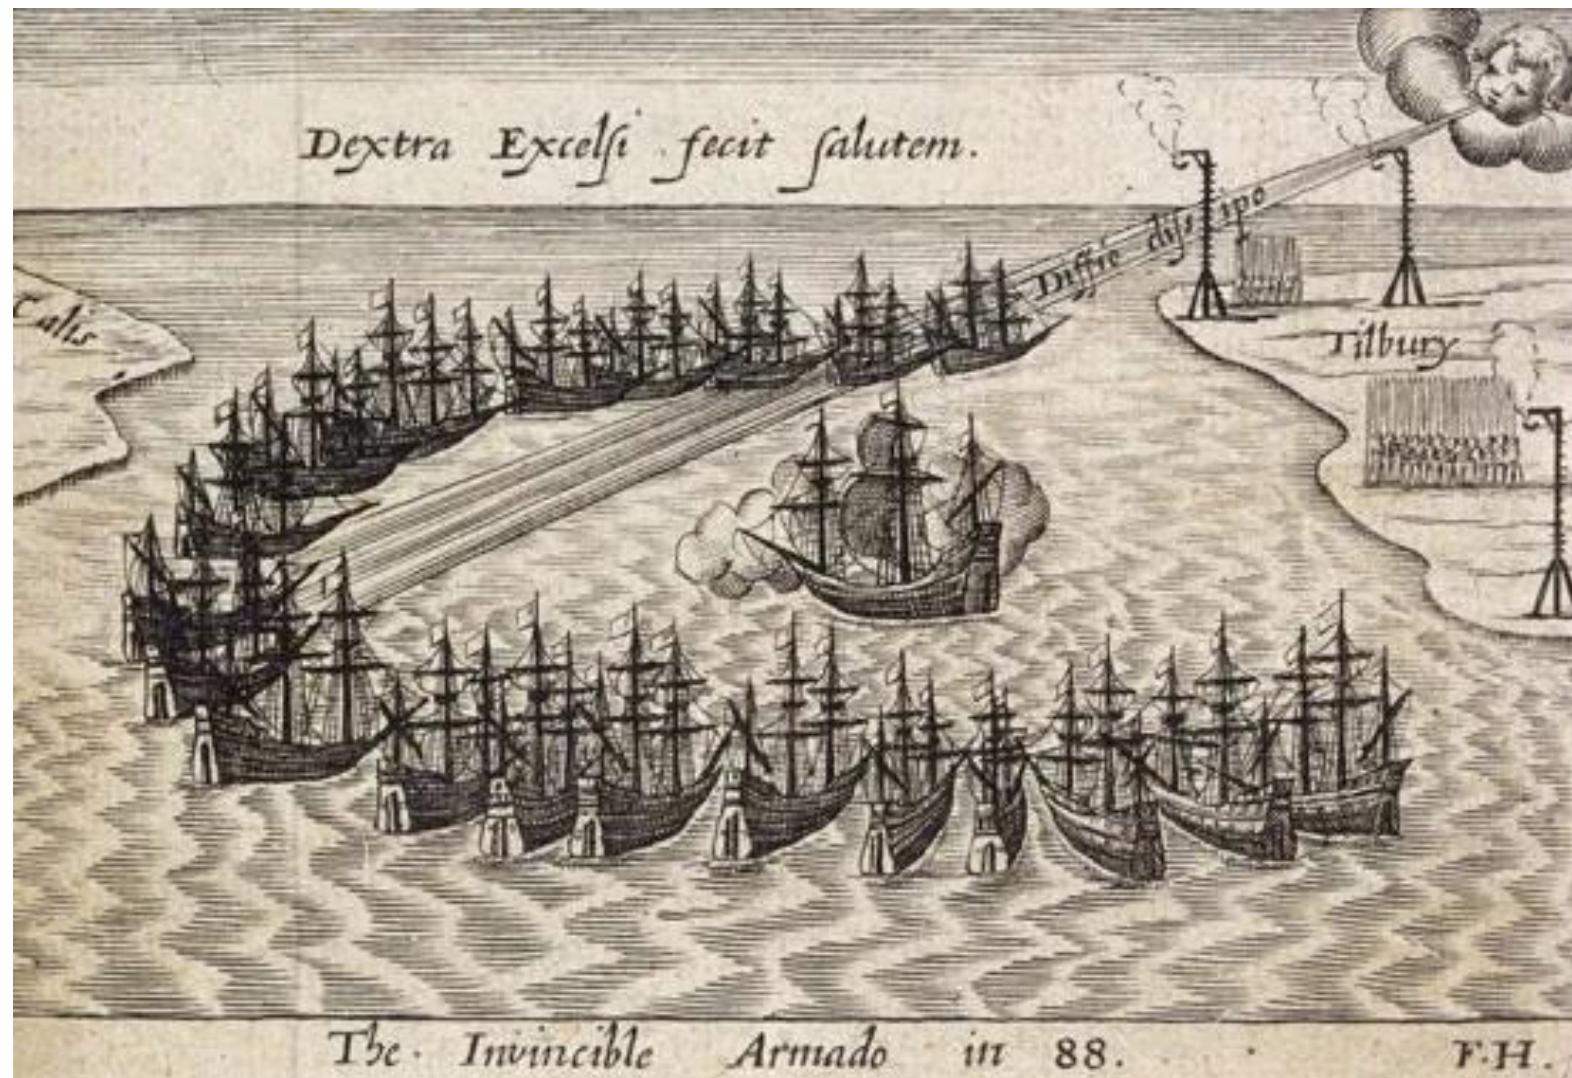


***Kubernetes*** is Greek for pilot or **helmsman**

# What is Kubernetes !?



- A helmsman
  - **maintains the course of the ship,**
  - **carries out the orders given by the captain**
  - **reports back the ship's heading.**
- Kubernetes
  - **steers applications**
  - **reports on their status**



# What is Kubernetes !?

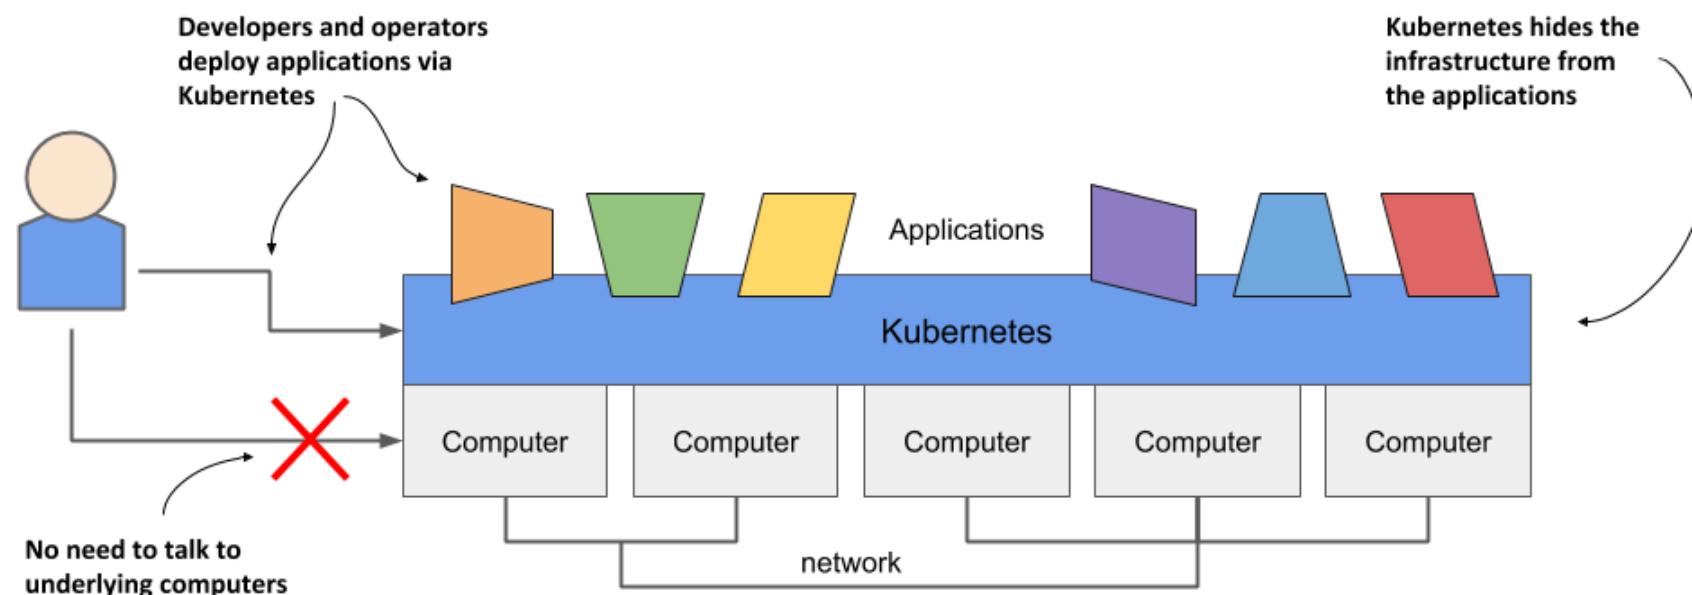


- Kubernetes enables **developers** to **deploy** their **applications themselves** and as often as they want
- The **focus for sysadmins** shifts from **supervising individual apps** to mostly **supervising and managing Kubernetes** and the **rest of the infrastructure**
  - while **Kubernetes** itself **takes care of the apps.**

# What is Kubernetes !?



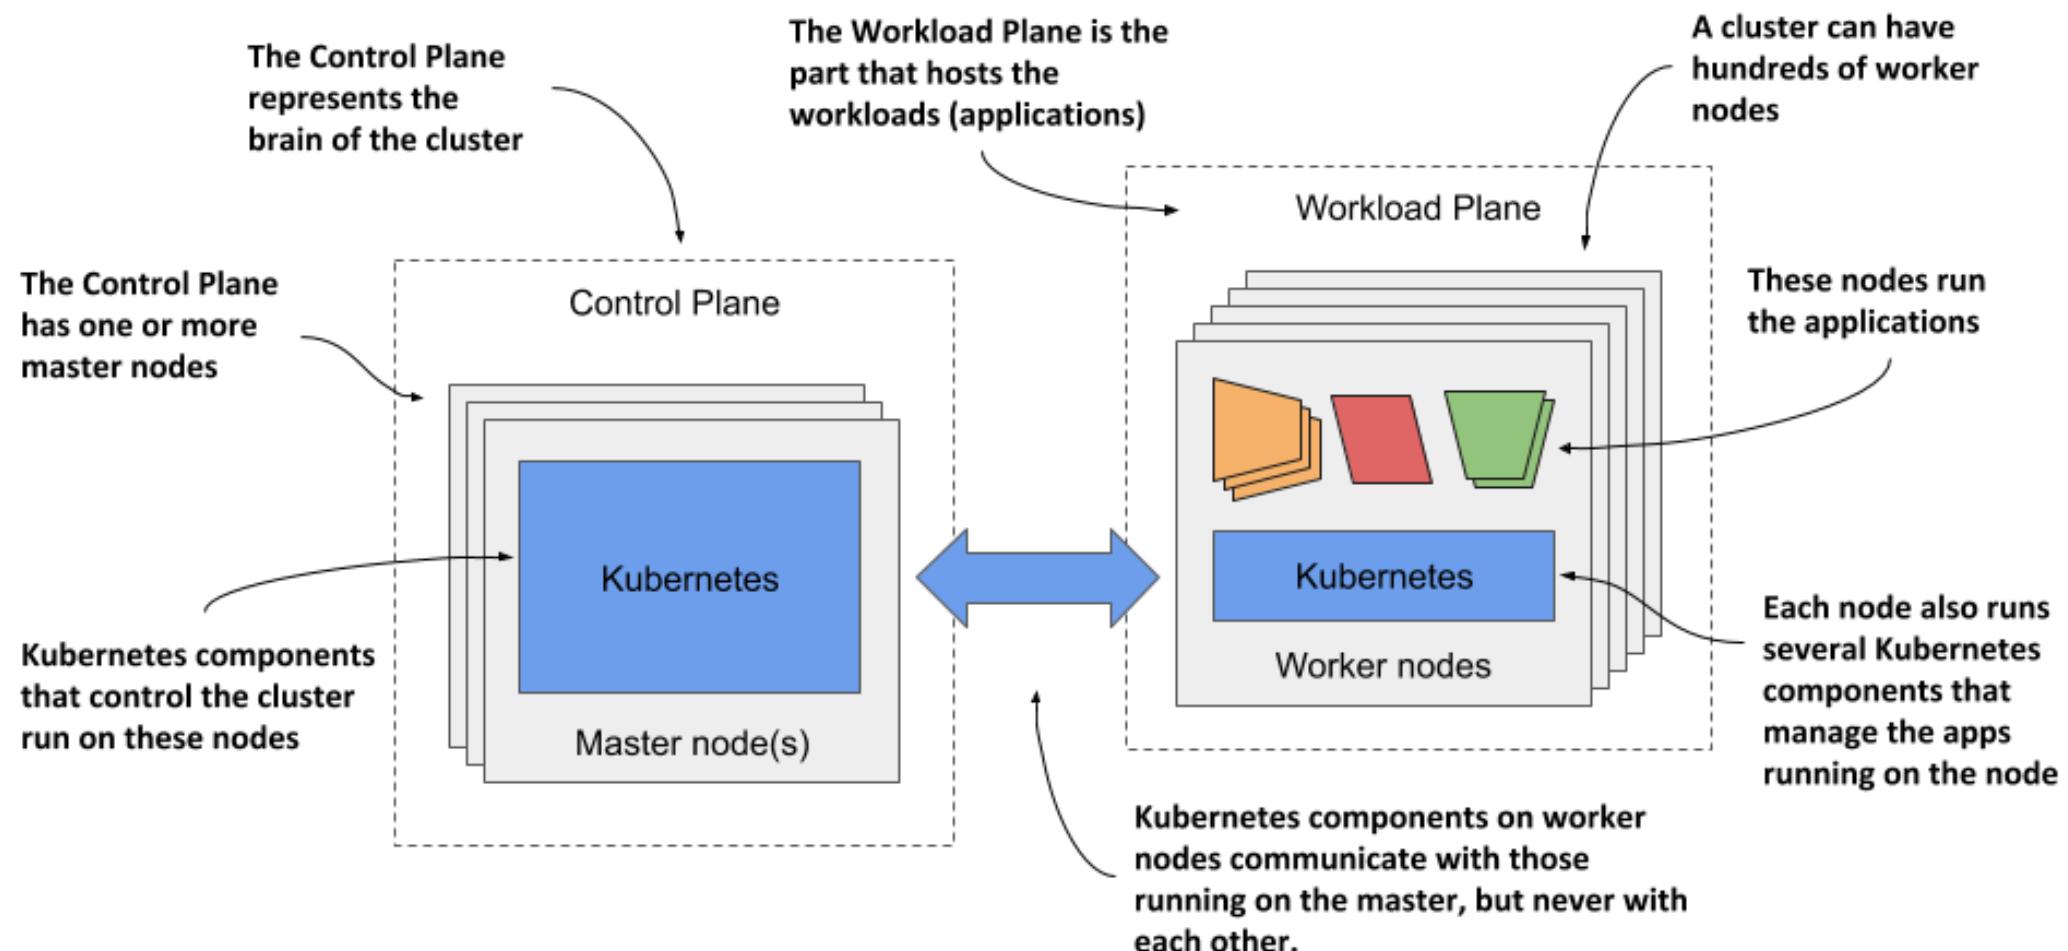
- It abstracts away the **underlying infrastructure**
- This **simplifies the development, deployment, and management** for both **development and the operations teams**.



# What is Kubernetes !?



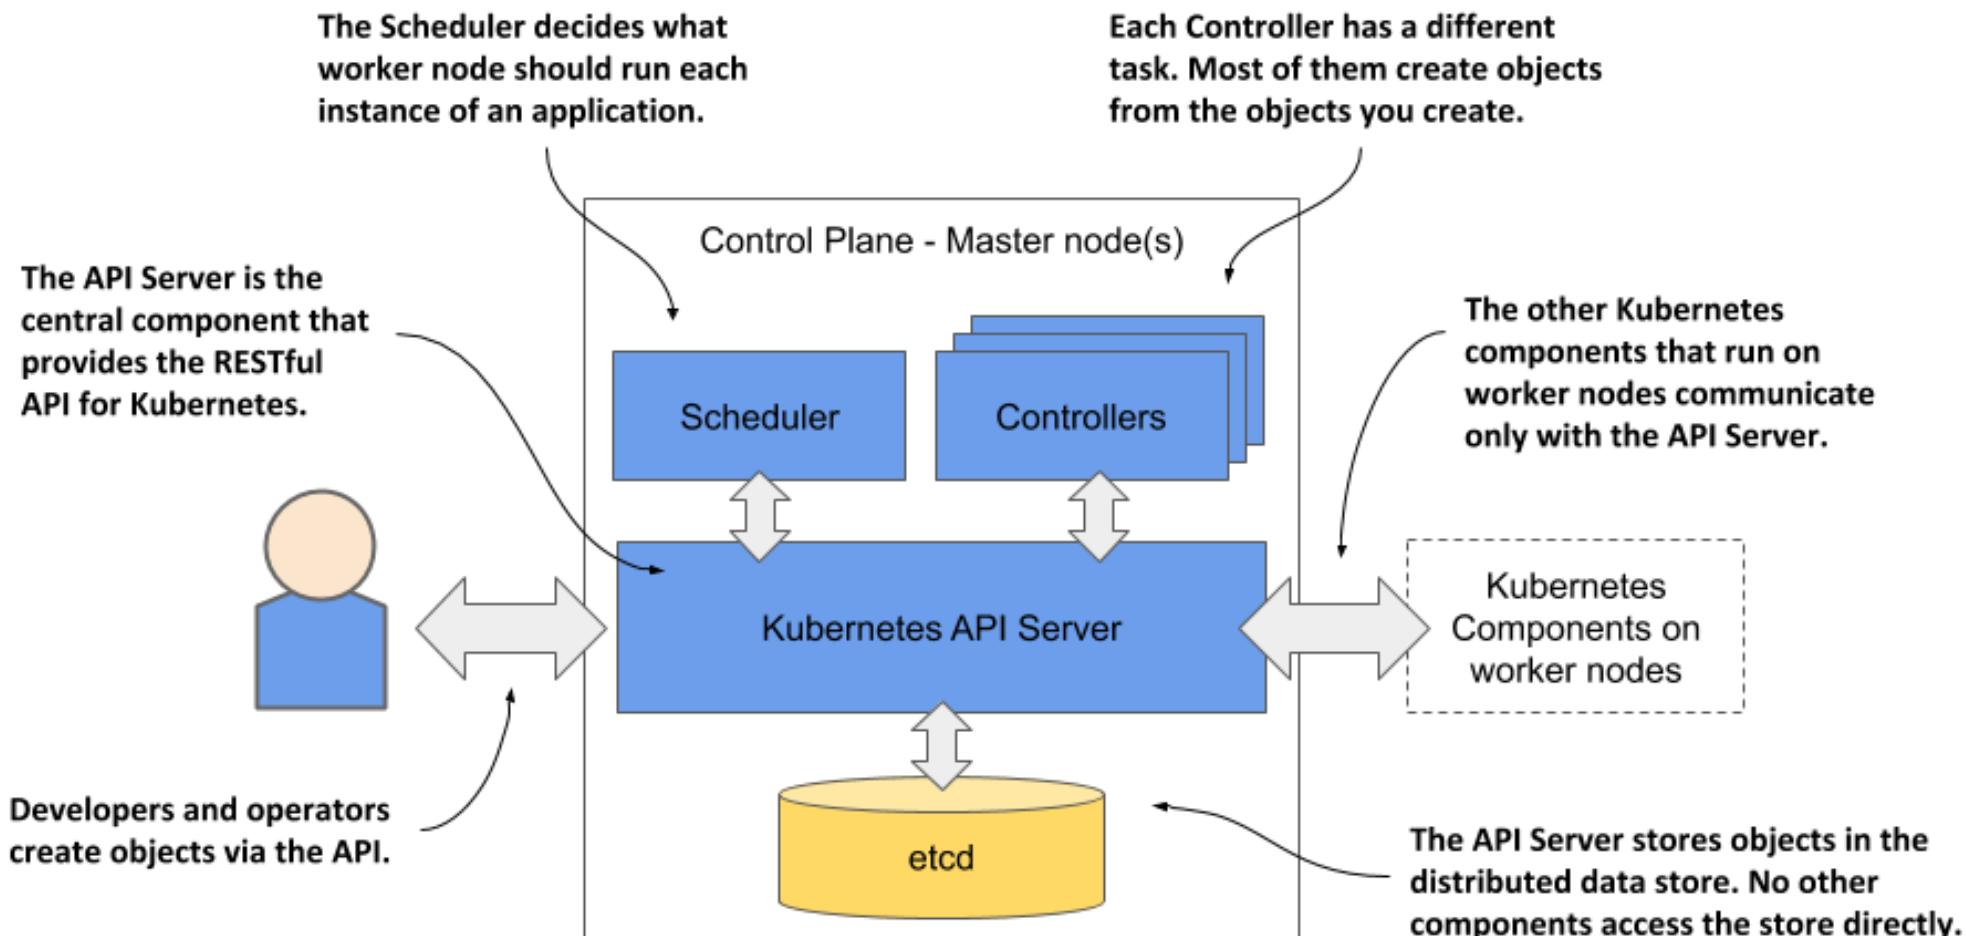
## Architecture overview



# What is Kubernetes !?



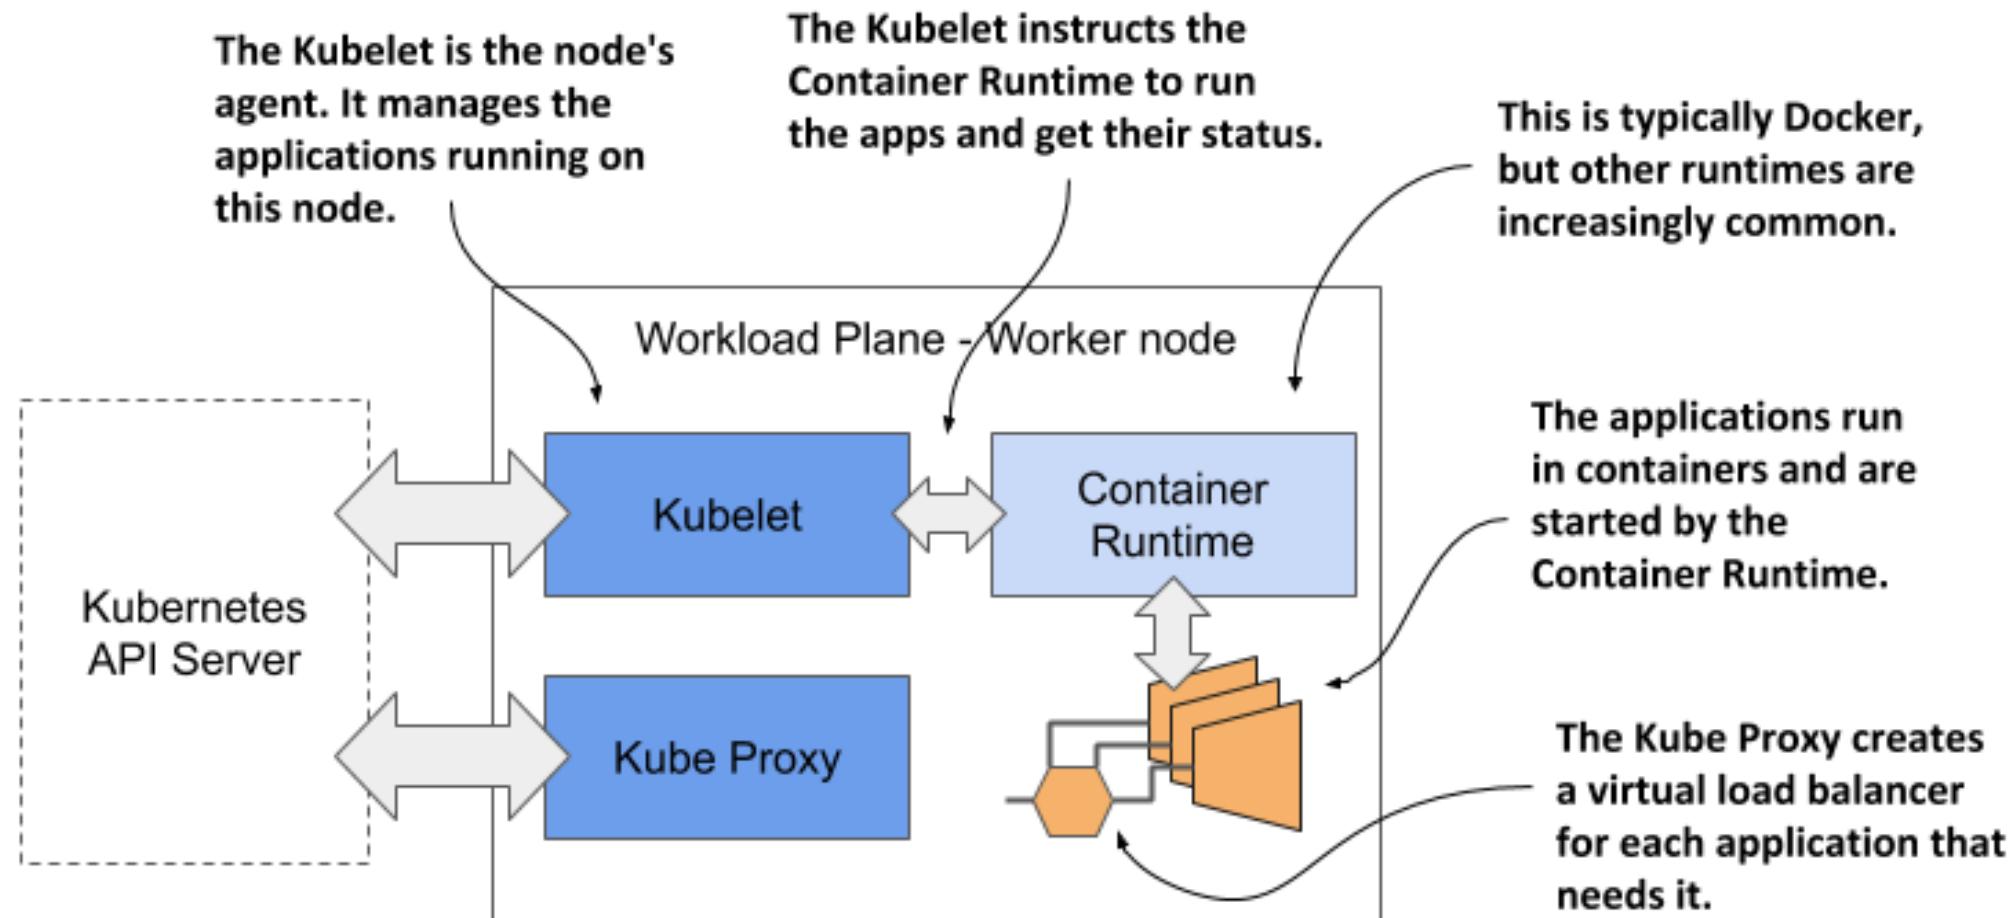
## The Control Plane



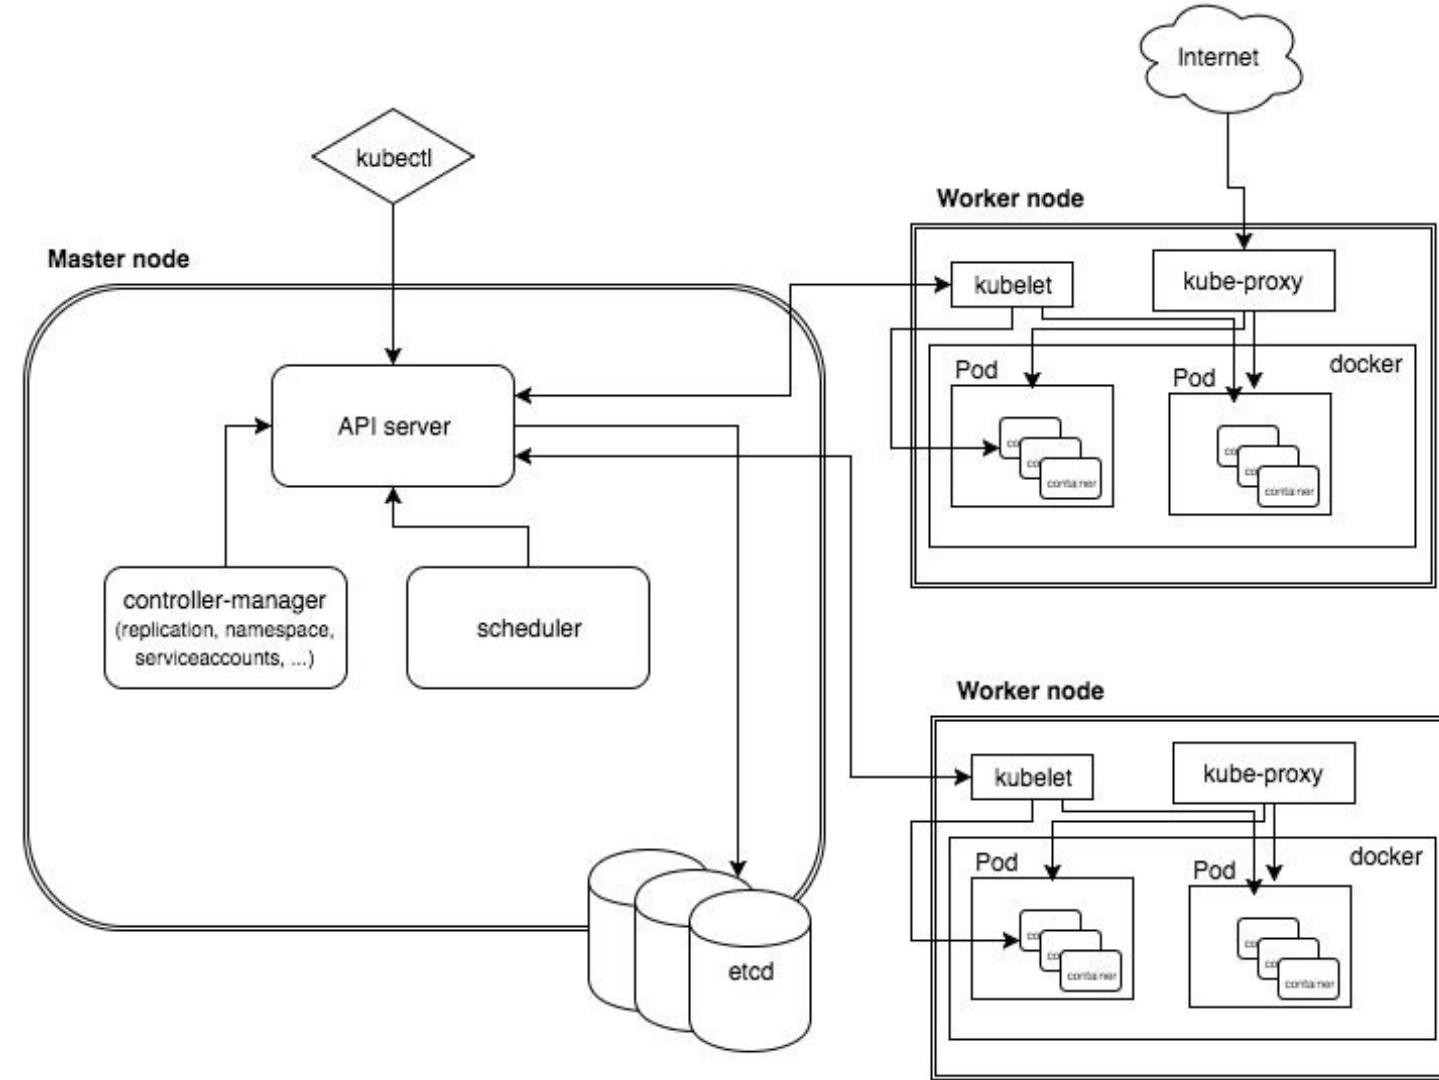
# What is Kubernetes !?



## The workload plane



# What is Kubernetes !?



# What is Kubernetes !?



## *The benefits*

*“Kubernetes does the things that the very **best system administrator would do: automation, failover, centralized logging, monitoring**. It takes what we’ve learned in the DevOps community and makes it the default, out of the box.*

—Kelsey Hightower

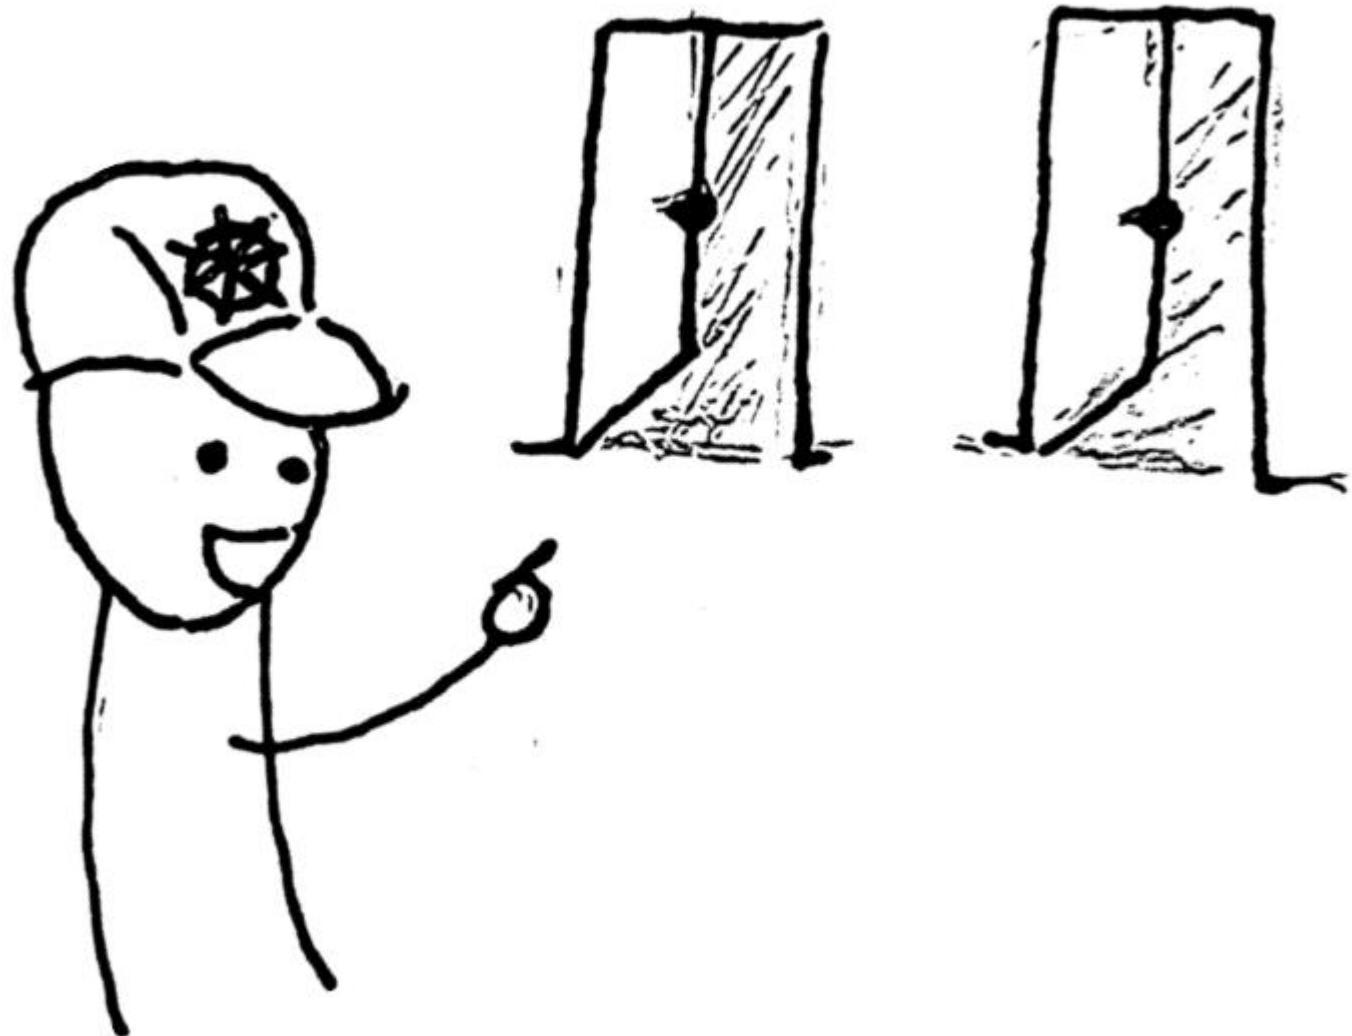
# What is Kubernetes !?



## *The benefits*

“*Everything Fails All the Time.*

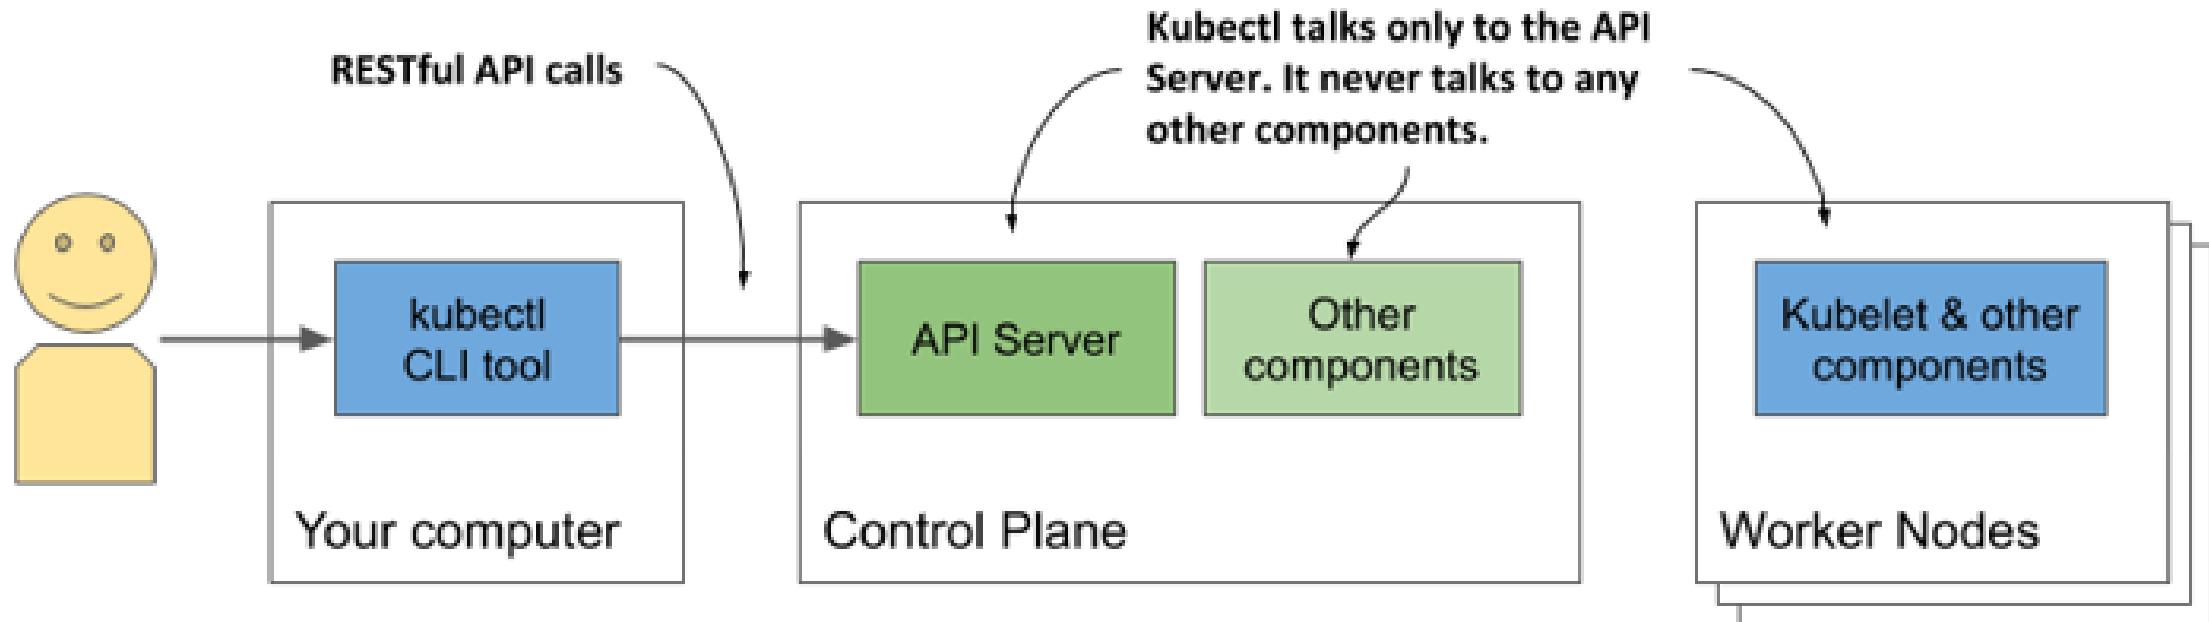
—Werner Vogels



# Interacting with Kubernetes



- **kubectl** is the **officially supported command-line tool** for accessing the Kubernetes API.
- It can be **installed on Linux, macOS, or Windows**.



# Interacting with Kubernetes



## *Completion for kubectl*

- kubectl command can also **output shell completion code** for **both the bash** and the **zsh shell**.
- It enables **tab completion** of **not only command names** but also the **object names**.
- With **tab completion**, things are much **easier**. Only **press TAB** after **typing the first few characters of each token**

# Interacting with Kubernetes



## *kubeconfig Configuration File*

- Such configuration files are referred to as **kubeconfig** files.
- Note that **kubeconfig** is a **generic way** to refer to **kubectl** configuration files and that it is **not the name of the config file**.
- **kubectl** uses such files to **store the information needed** for us to **choose a cluster** and **communicate** with its **API** server.

# Interacting with Kubernetes



## *kubeconfig Configuration File*

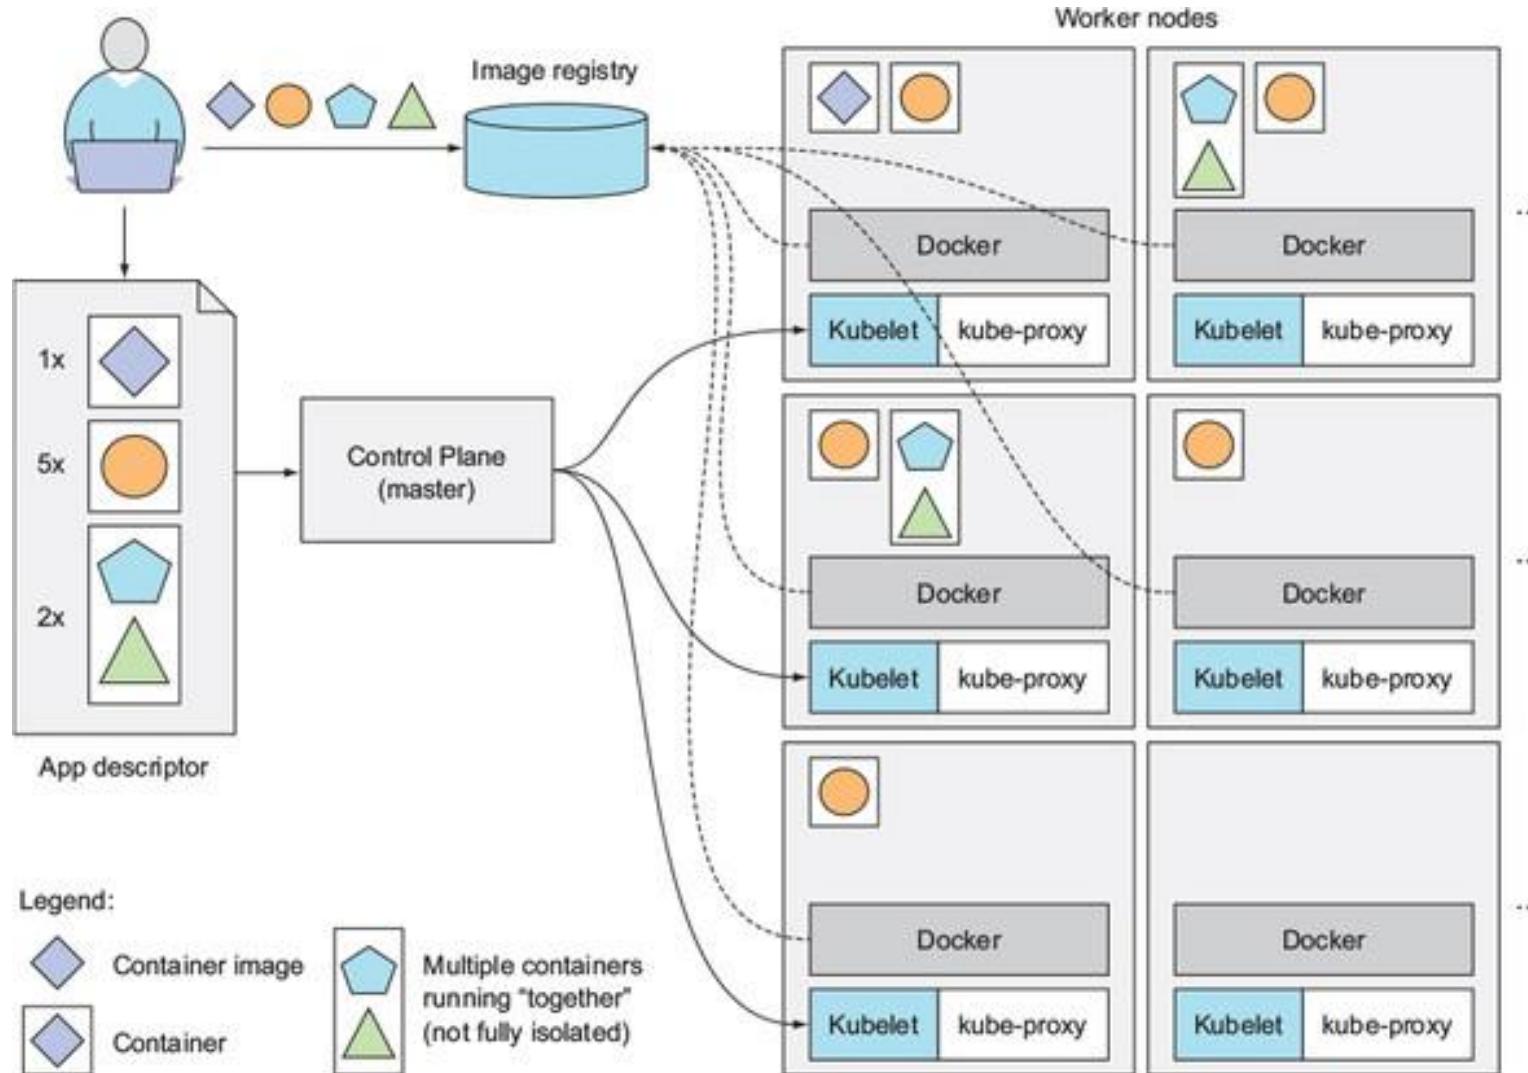
- By **default**, kubectl looks for the file in the **\$HOME/.kube** directory.
- In most scenarios, specify a **KUBECONFIG** environment variable or use the **--kubeconfig** flag to specify the **kubeconfig** files.

# Running an application in Kubernetes

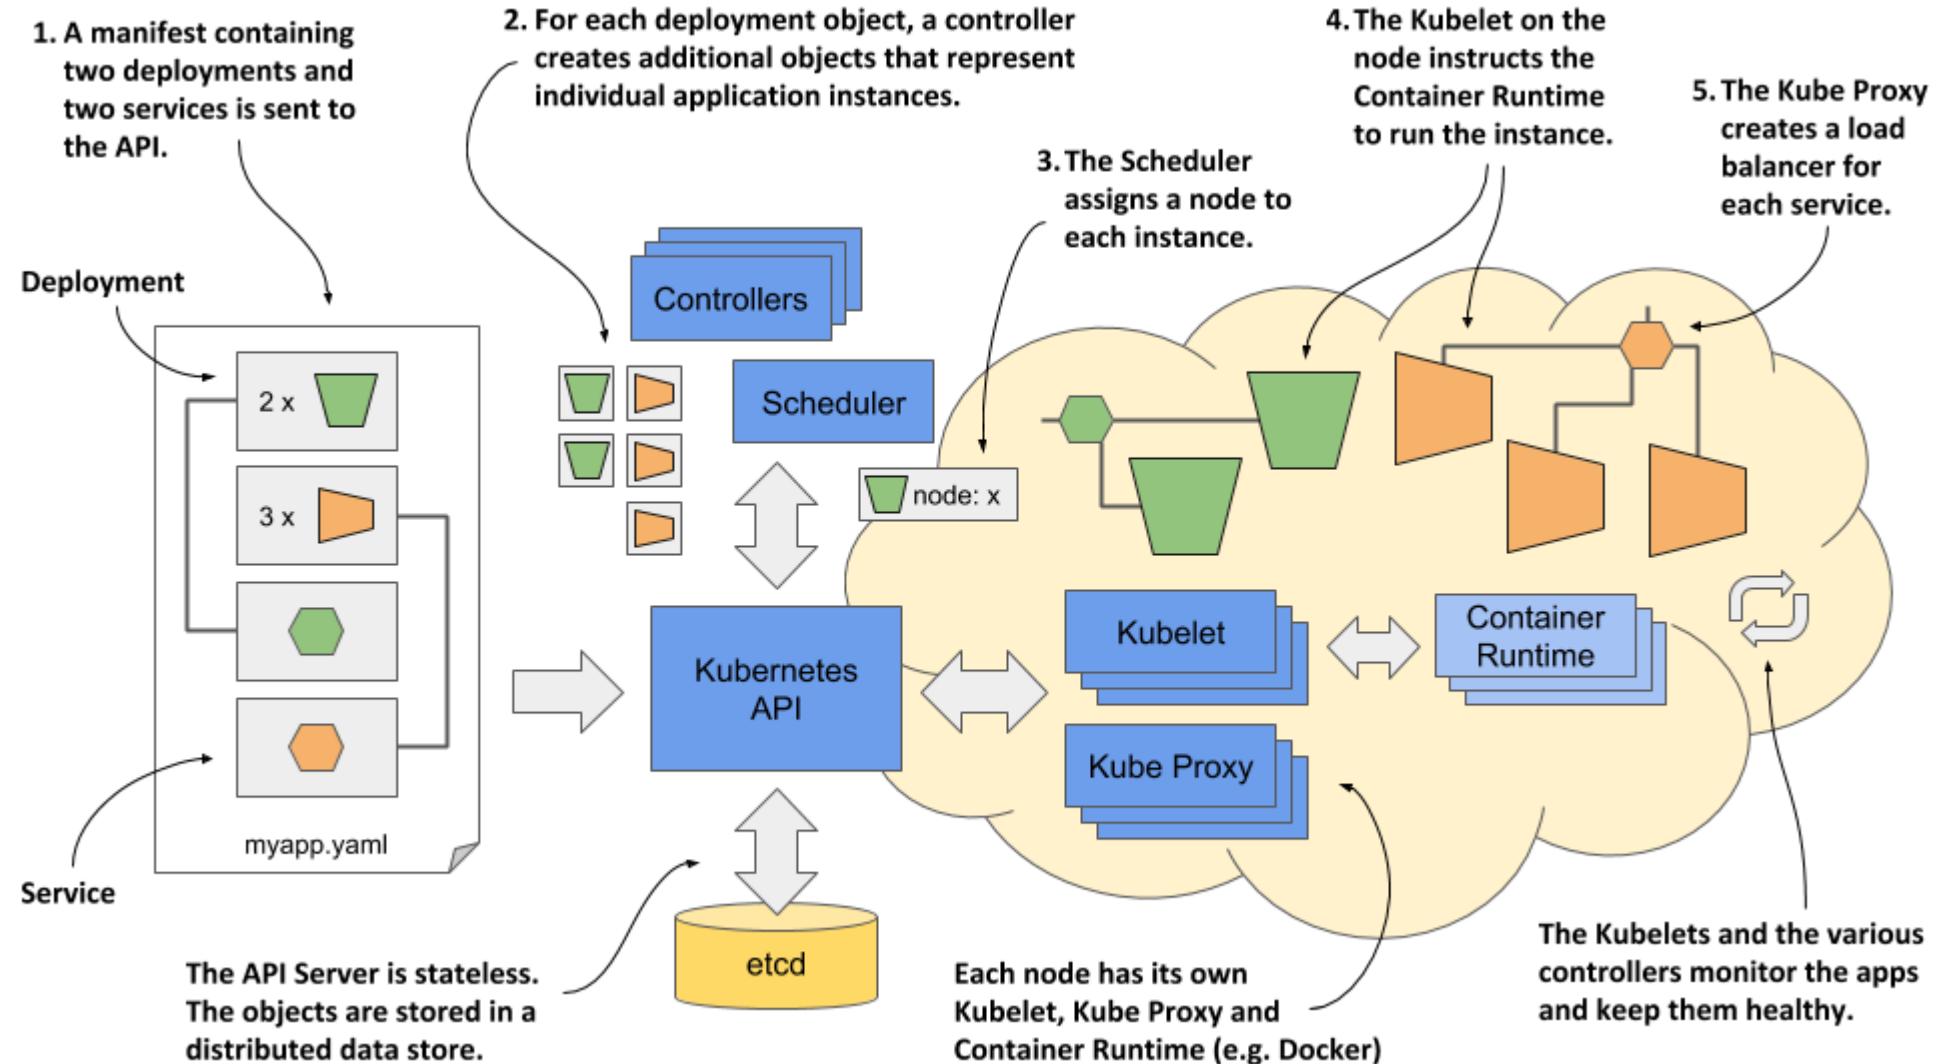


- **Everything** in Kubernetes is **represented** by an **object**.
- These **objects** can be **created** and **retrieved** via the **Kubernetes API**.
- These **objects** are usually **defined** in one or more **manifest** files in either **YAML** or **JSON** format.

# Running an application in Kubernetes



# Running an application in Kubernetes



# Running an application in Kubernetes



- Kubernetes **continuously makes sure that the deployed state of the application always matches the provided description.**
- Kubernetes also **reports the status of the application by updating the object that represents the application instance.**
- The other **controllers monitor these objects and ensure that applications are moved to healthy nodes if their nodes fail.**

# Running an application in Kubernetes



- Kubernetes **allows clients** to easily **find containers** that **provide a specific service**
- Kubernetes will **expose all** of them at a **single static IP** address and **expose** that **address** to all **applications** running in the **cluster**.
- This is done **through environment variables**, but **clients** can also **look up** the **service IP** through good old **DNS**.

# Pods



- A **container** is not the **smallest unit of deployment** in **Kubernetes**.
- A **POD** is the **smallest object**, that you **can create** in **kubernetes**.



# Pods



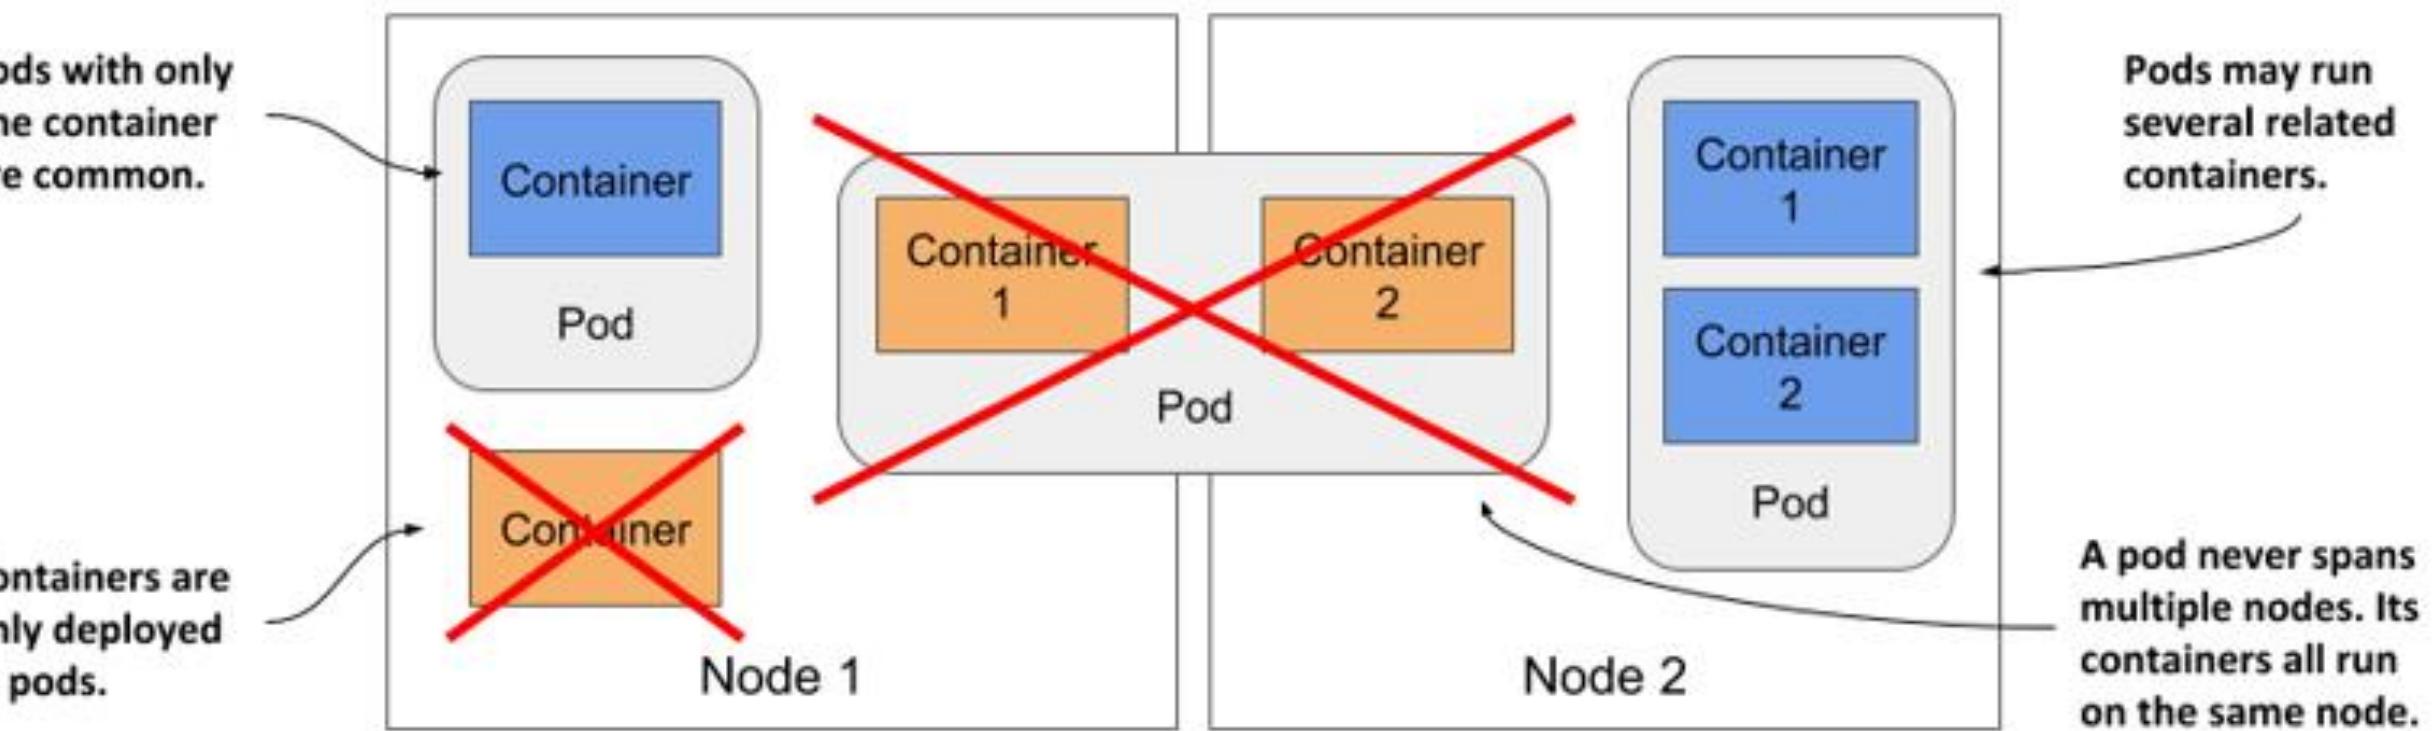
CHUSANALPSEEN

# Pods

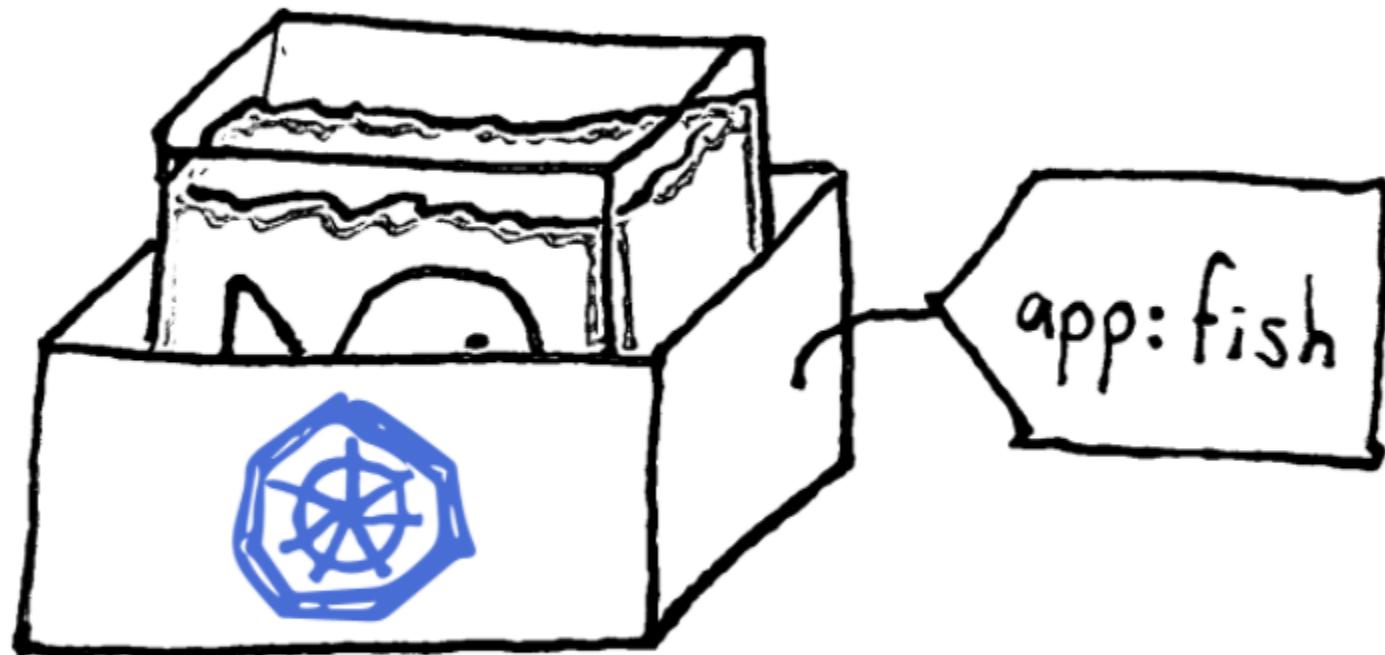


Pods with only one container are common.

Containers are only deployed in pods.

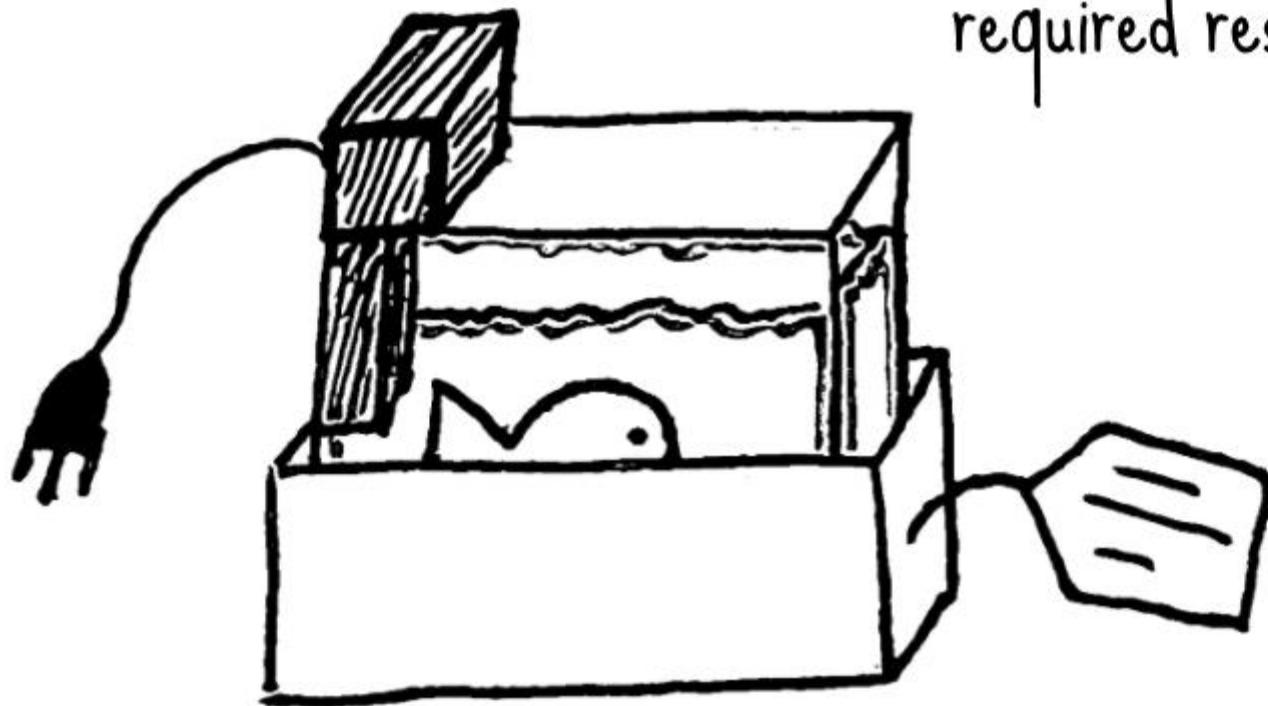


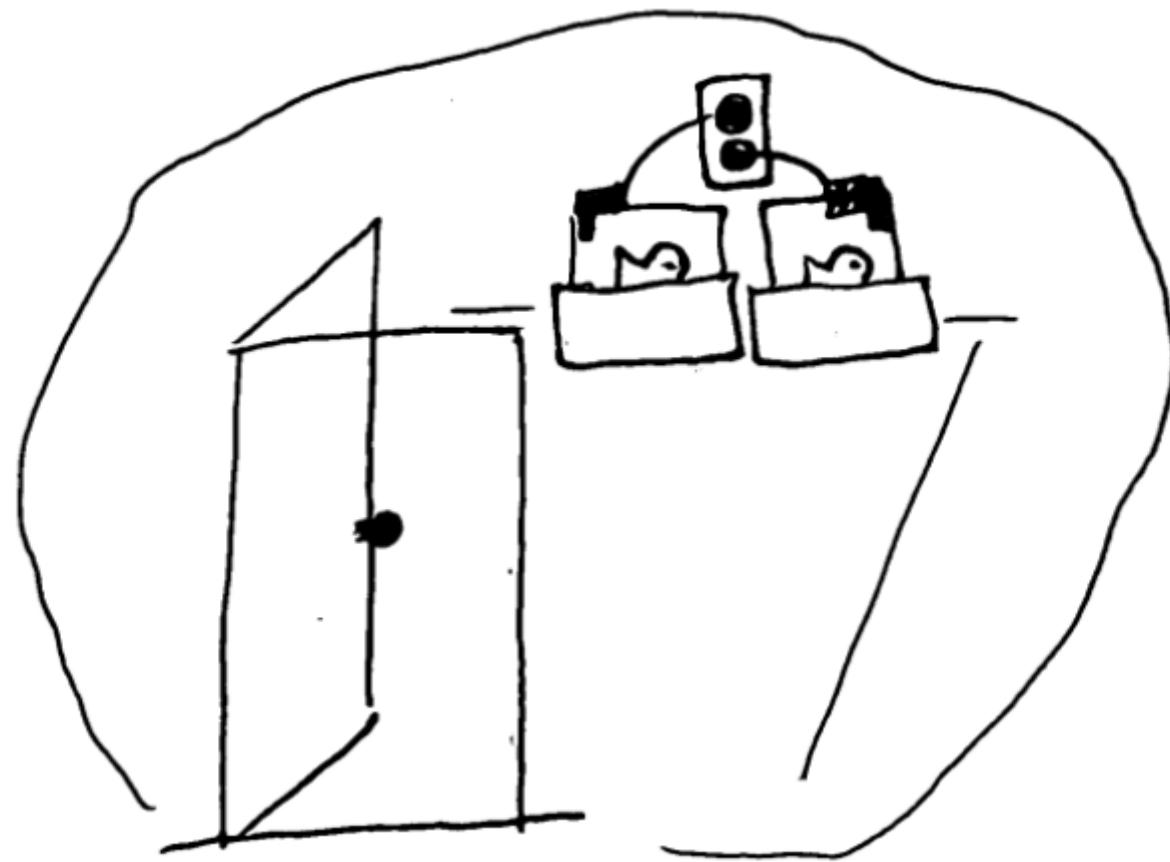
pod

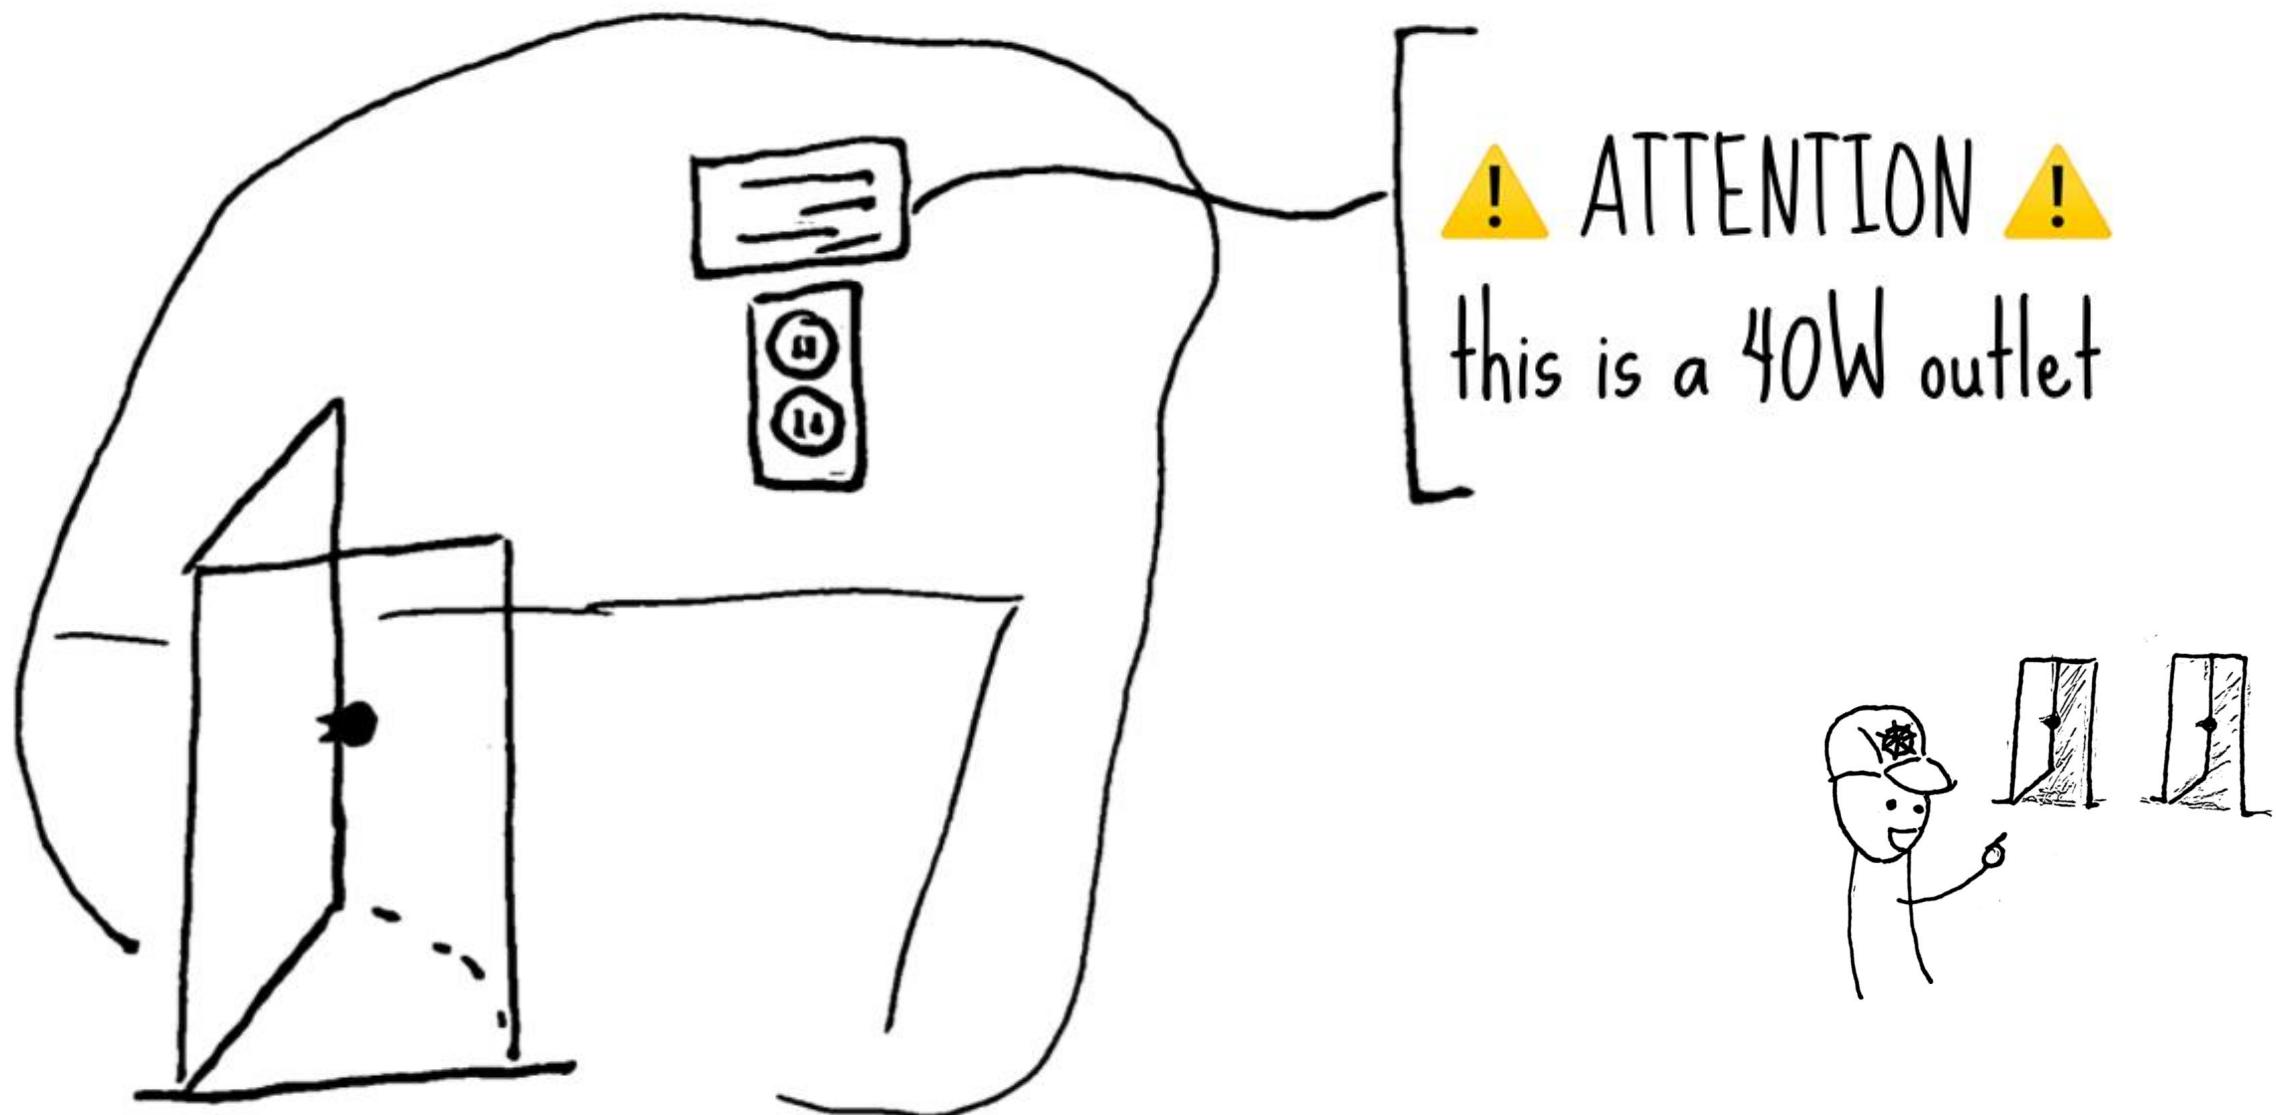


water filter:

required resources: 60W





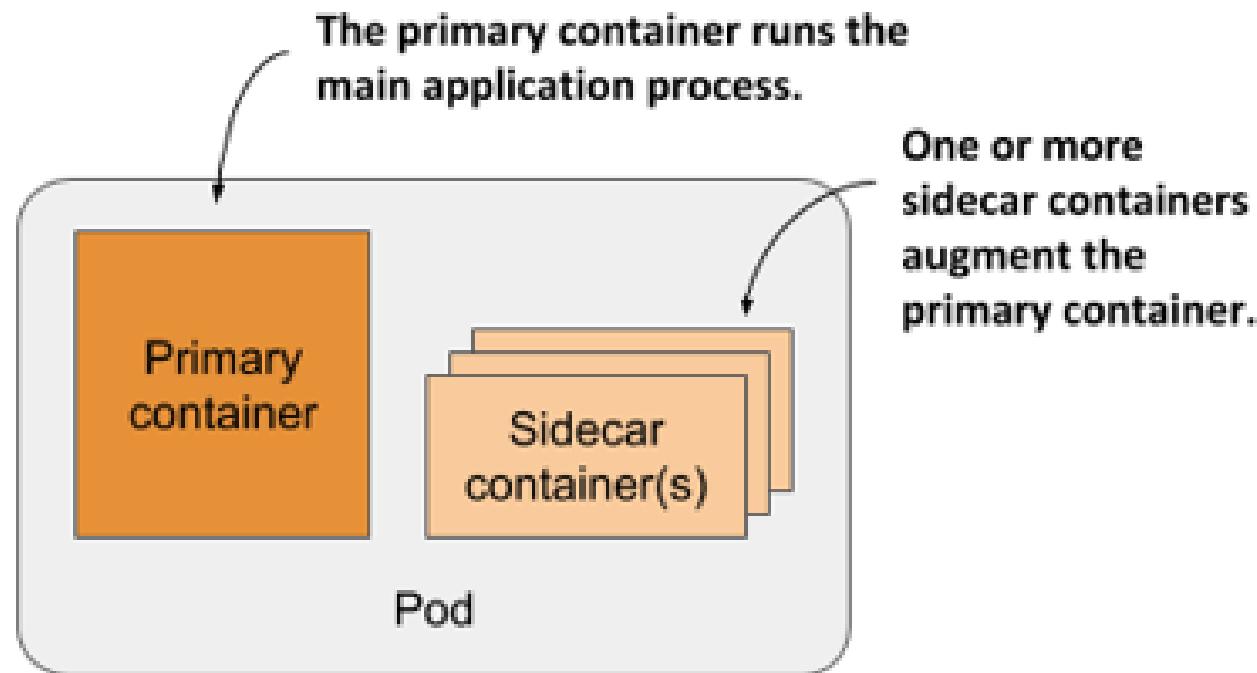


# Pods



## Multiple containers in a pod

- Multiple containers into a single pod is used for applications that consists of one main process and one or more complementary processes



# Pods



## Pod creation with cli

- The simplest way to create a Pod is via the imperative kubectl run command:

***\$ kubectl run nginx --image nginx***

- To see the status of the pod

***\$ kubectl get pods***

- To delete the pod

***\$ kubectl delete pods/nginx***

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

This manifest uses the v1 API version to define the object

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

The object specified in this manifest is a  
**Pod**

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

The name of the pod

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

- The name of the container
- Container image to create the container from
- The port the app is listening on

# Pods



## Pod creation with manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: demo-app
    image: luksa/kubia:1.0
  ports:
  - containerPort: 8080
```

Use the **kubectl apply** command to launch a single instance of the pod:

**\$ *kubectl create -f demo-pod.yaml***

# Pods



## TIP

- To create a pod **manifest from scratch**, use the following command to **create** the file and **then edit** it to add **more fields**:

```
$ kubectl run demo-app --image=luksa/kubia:1.0 --dry-run=client -o yaml > mypod.yaml
```

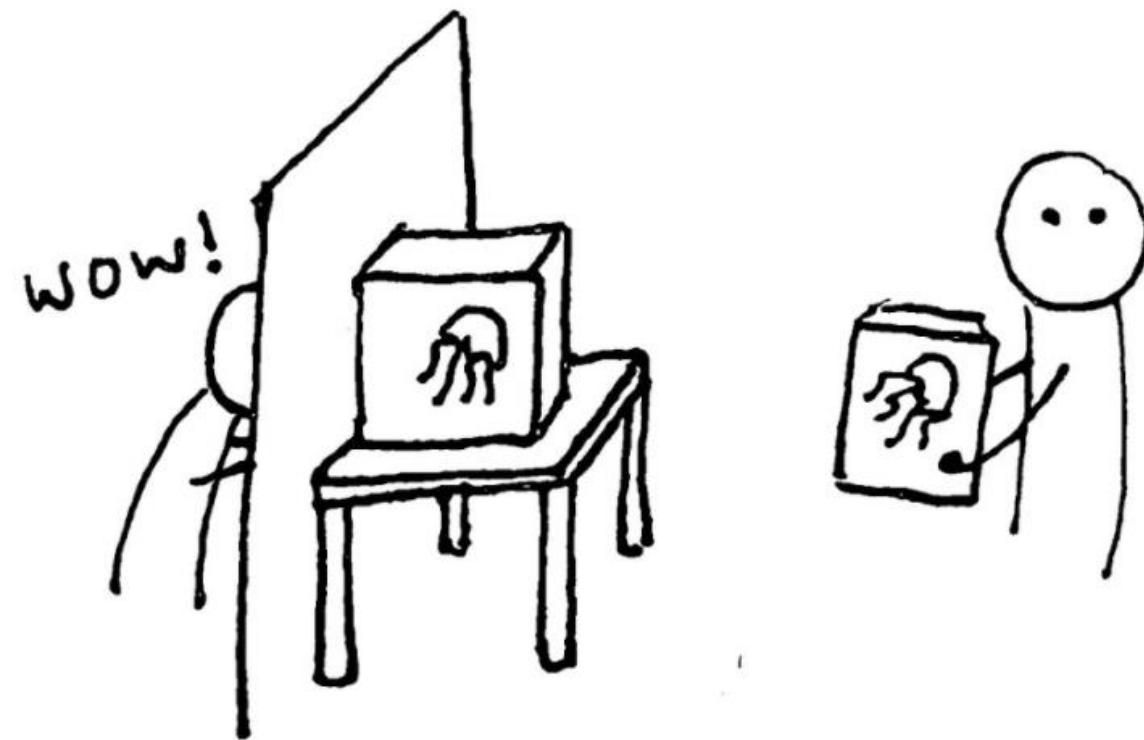
- The **--dry-run=client** flag tells kubectl to output the definition instead of actually creating the object via the API.



- Now we know how to:
  - **package** our **app components** into containers,
  - **group them** into pods,
- Let's see how we can **manage** them and **update** them.

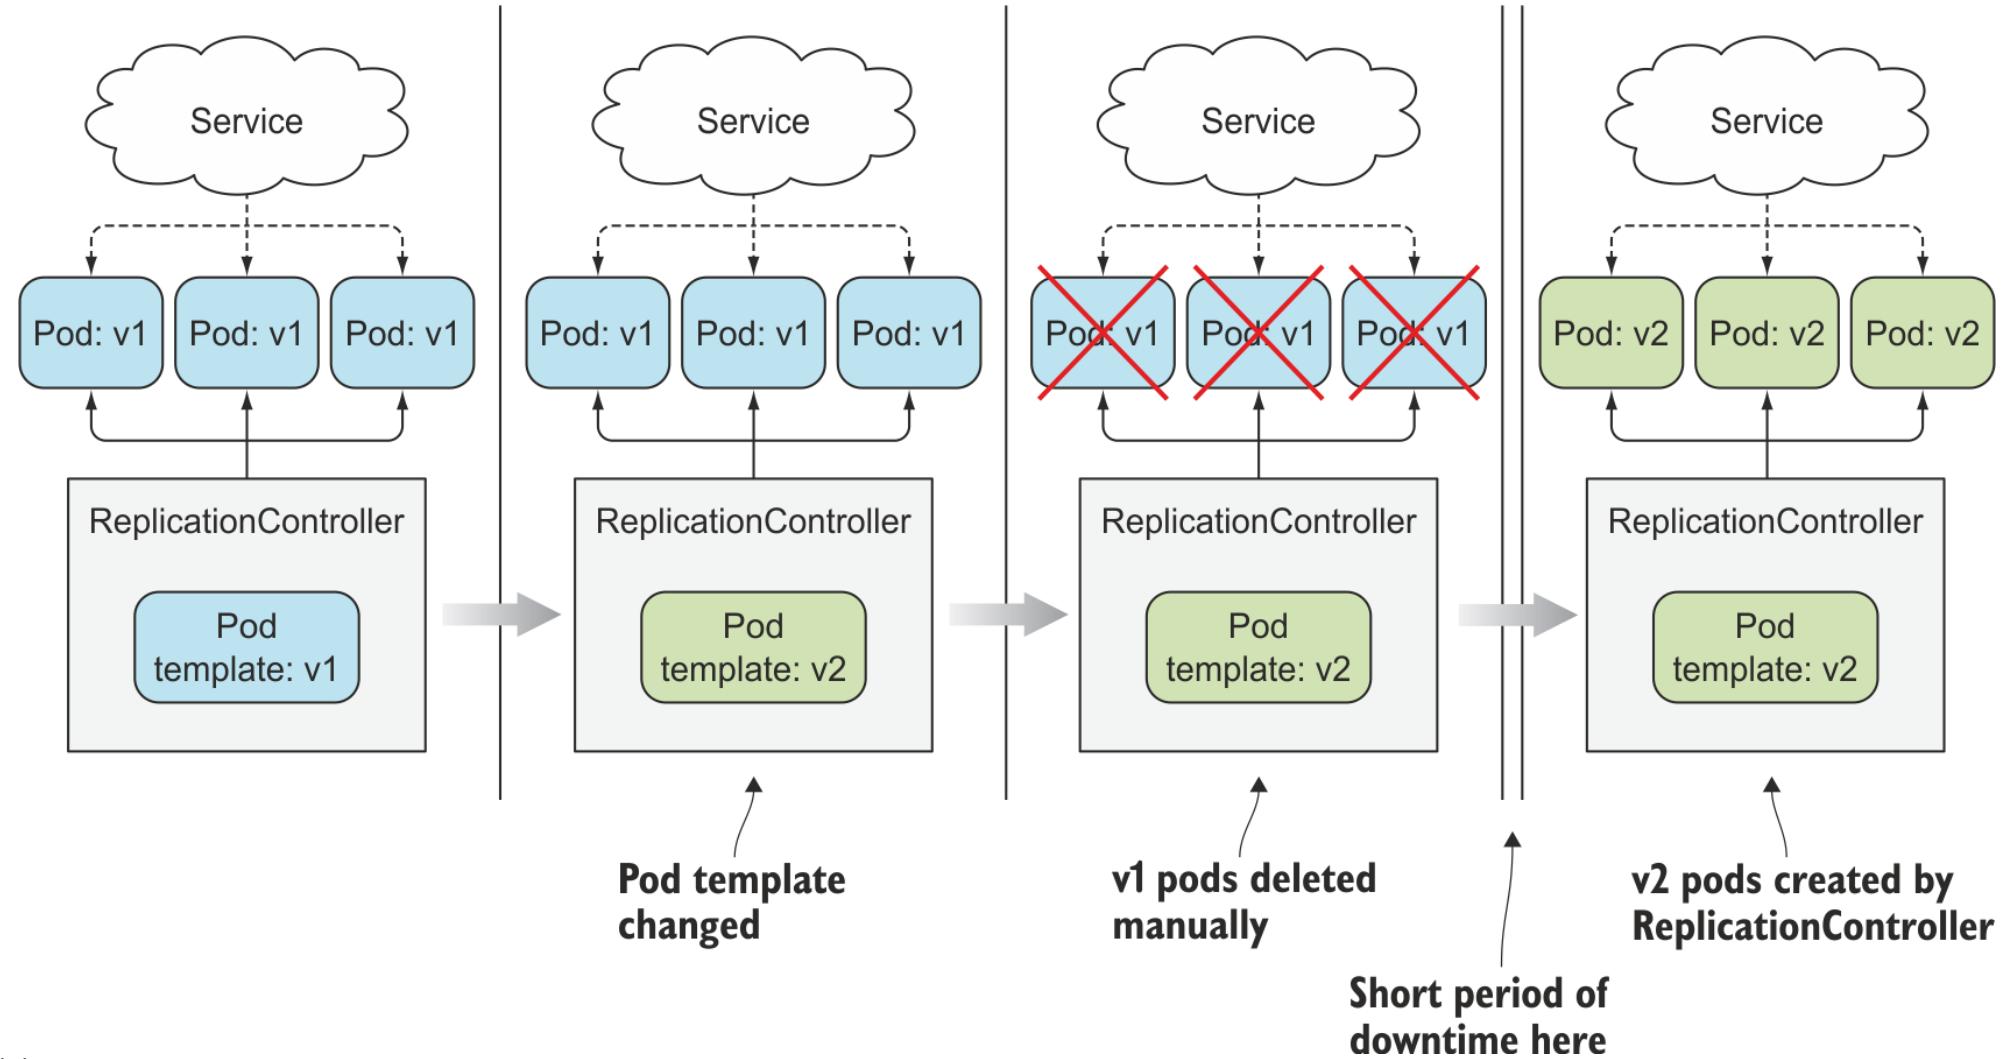
replica set

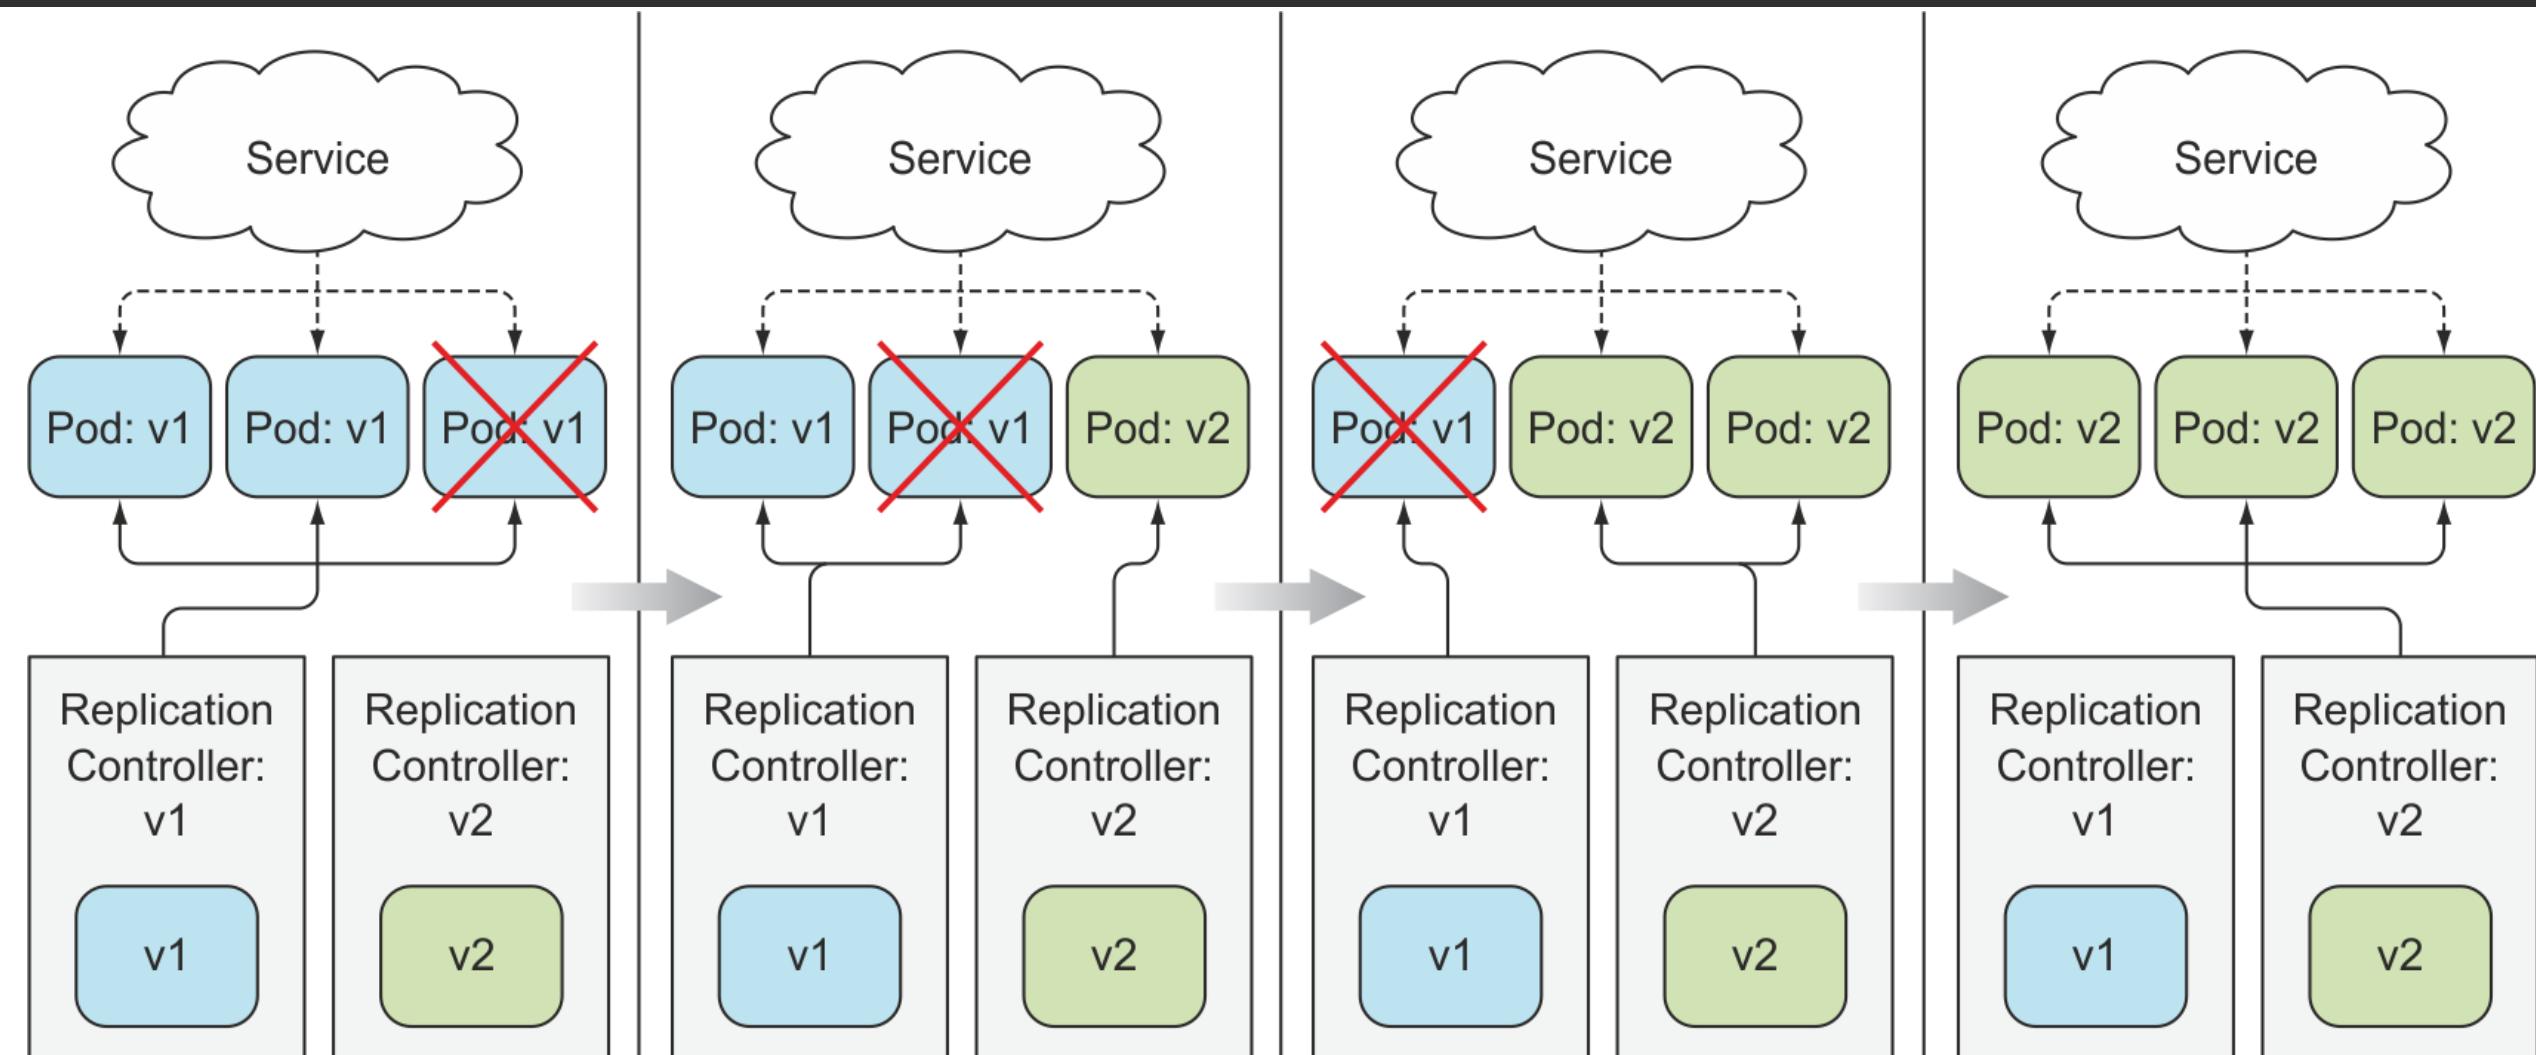






- There are two ways of **updating** pods





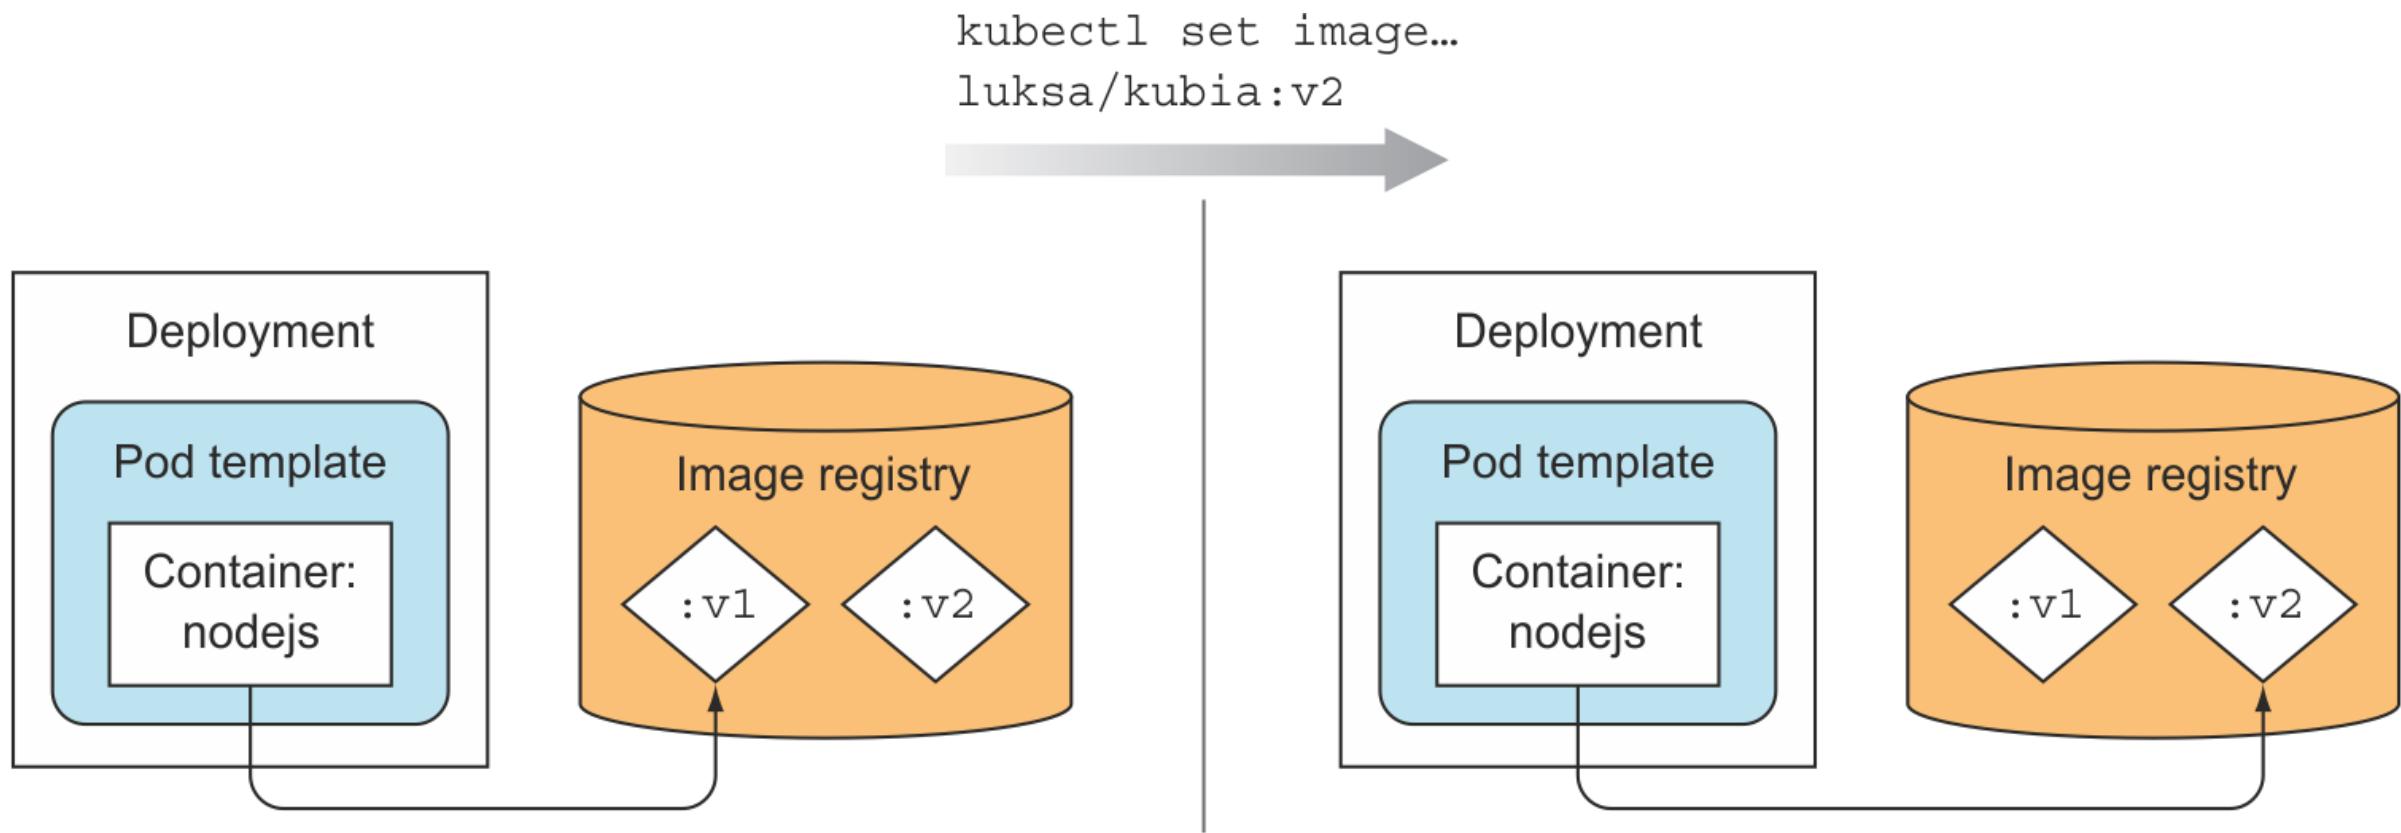
# Deployments



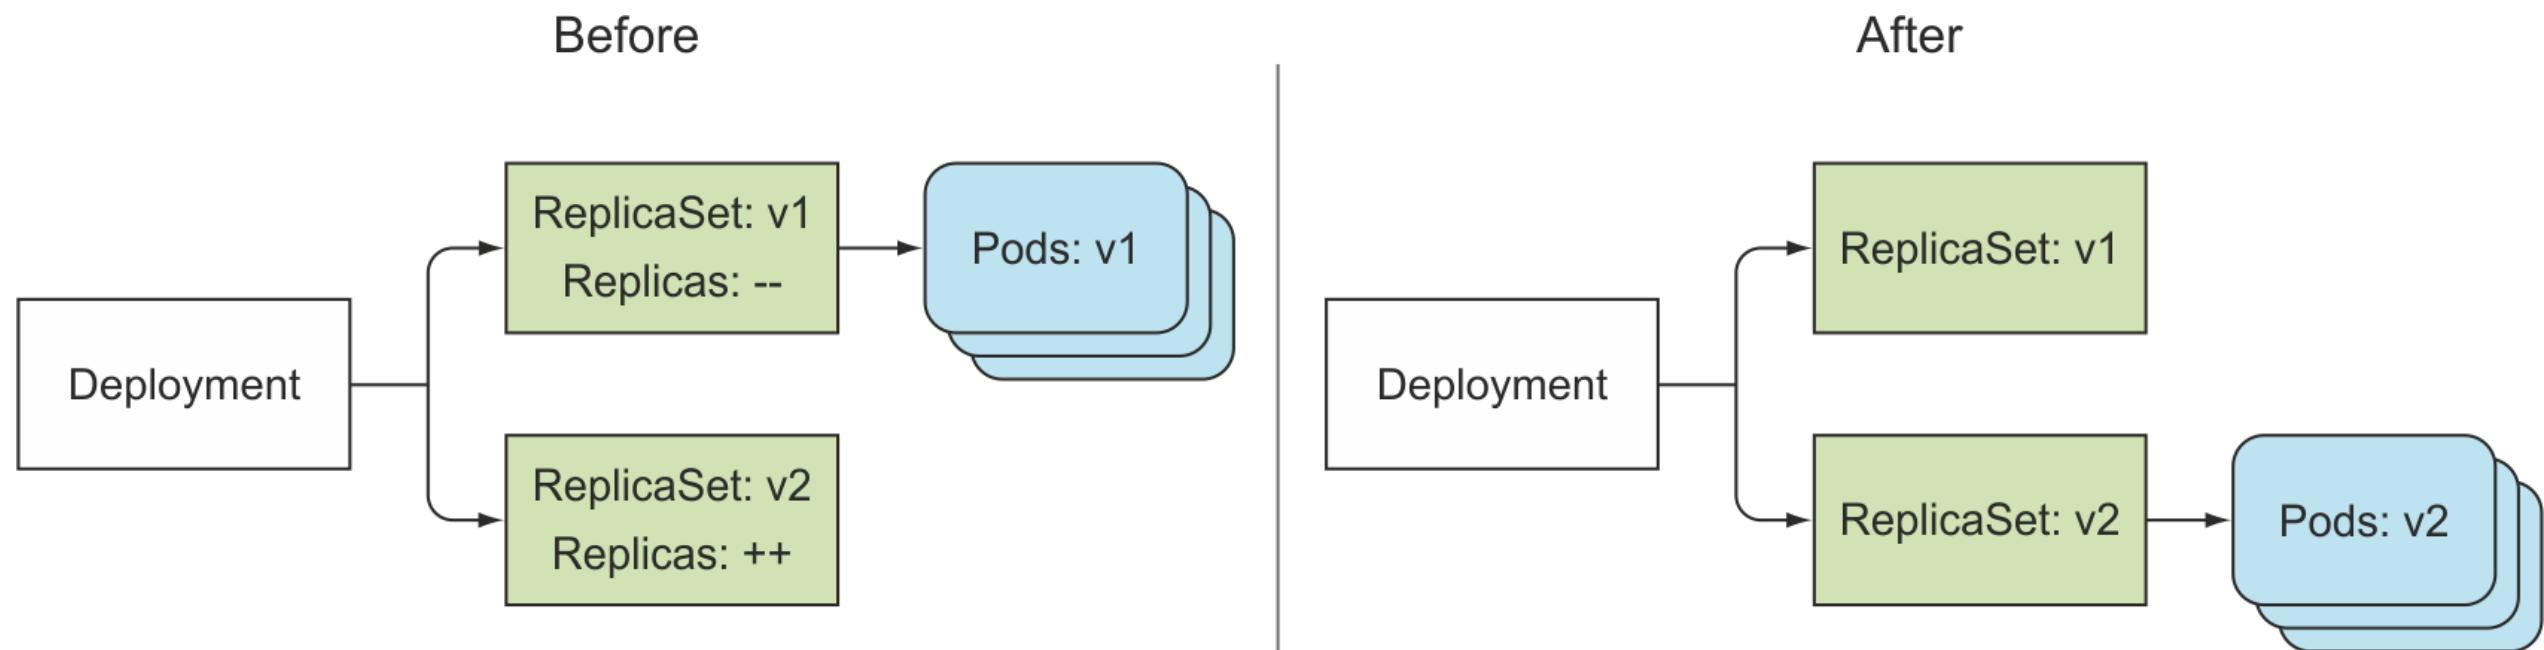
- A Deployment is a **higher-level resource** meant for **deploying applications** and **updating**
- When using a Deployment, the **actual pods** are **created and managed** by the Deployment's **ReplicaSets**, **not** by the **Deployment** directly



# Deployments

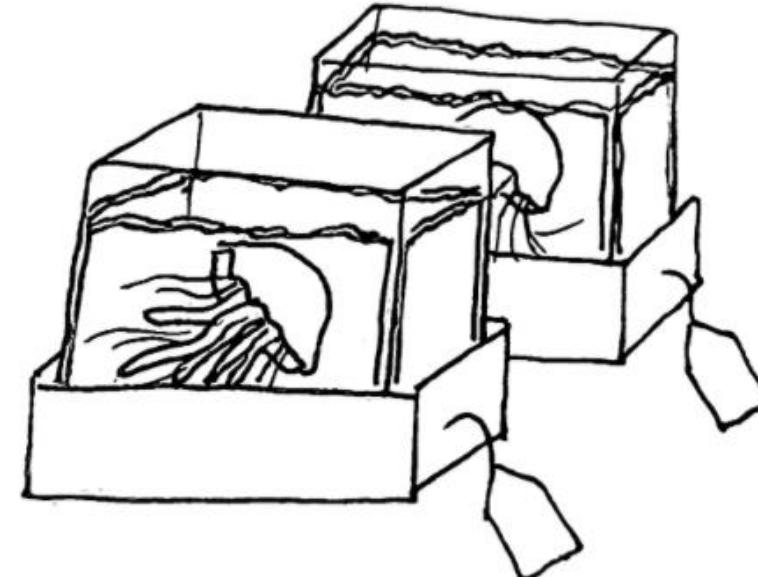
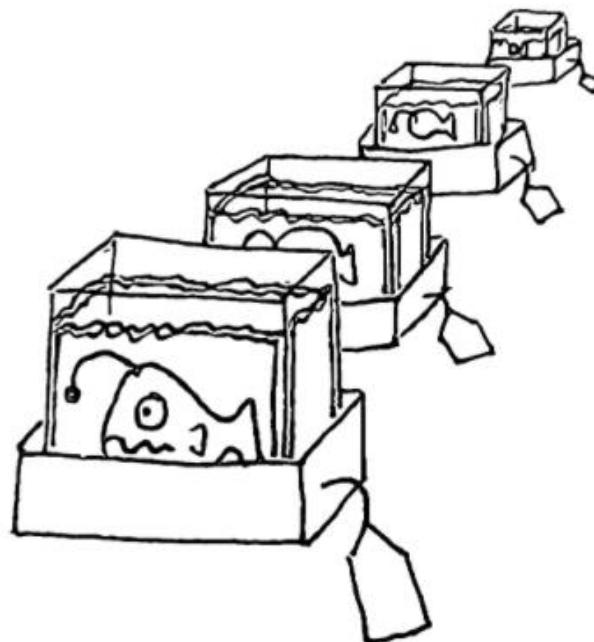
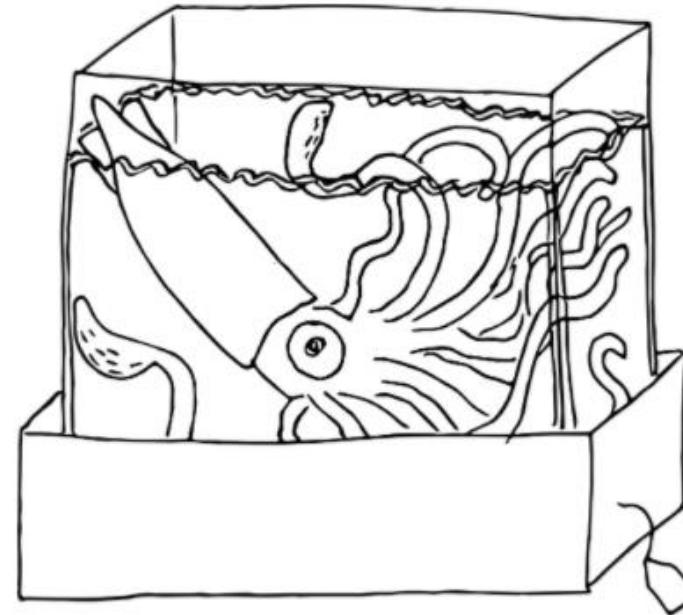


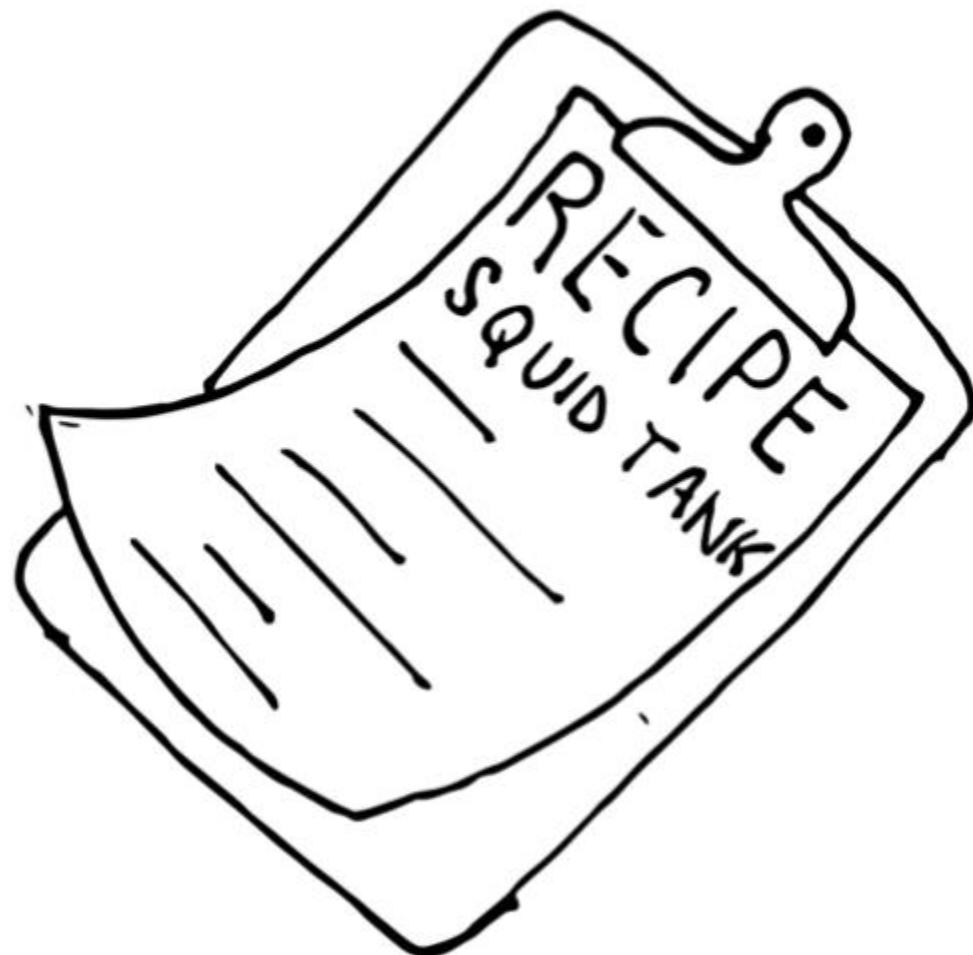
# Deployments



# deployment

## the Deep Sea exhibit





# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
  selector:
    matchLabels:
      app: kubia
```

\$ **kubectl create -f kubia-deployment-v1.yaml --record**

- **--record** **records** the command in the **revision history**

\$ **kubectl get deployment**  
\$ **kubectl describe deployment**  
\$ **kubectl rollout status deployment kubia**  
\$ **kubectl get po**  
\$ **kubectl get replicases**

# Deployments



## Rolling back - Simulation

```
$ kubectl set image deployment kubia nodejs=luksa/kubia:v3  
$ kubectl rollout status deployment kubia  
$ watch kubectl get po  
$ kubectl rollout undo deployment kubia
```

# Deployments



## **Rolling back** – Simulation

- Rolling back a rollout is **possible** because Deployments keep a **revision history**.
- The **history** is stored in the **underlying ReplicaSets**.
- When a **rollout completes**, the **old ReplicaSet isn't deleted**, and this **enables rolling back** to any revision

# Deployments



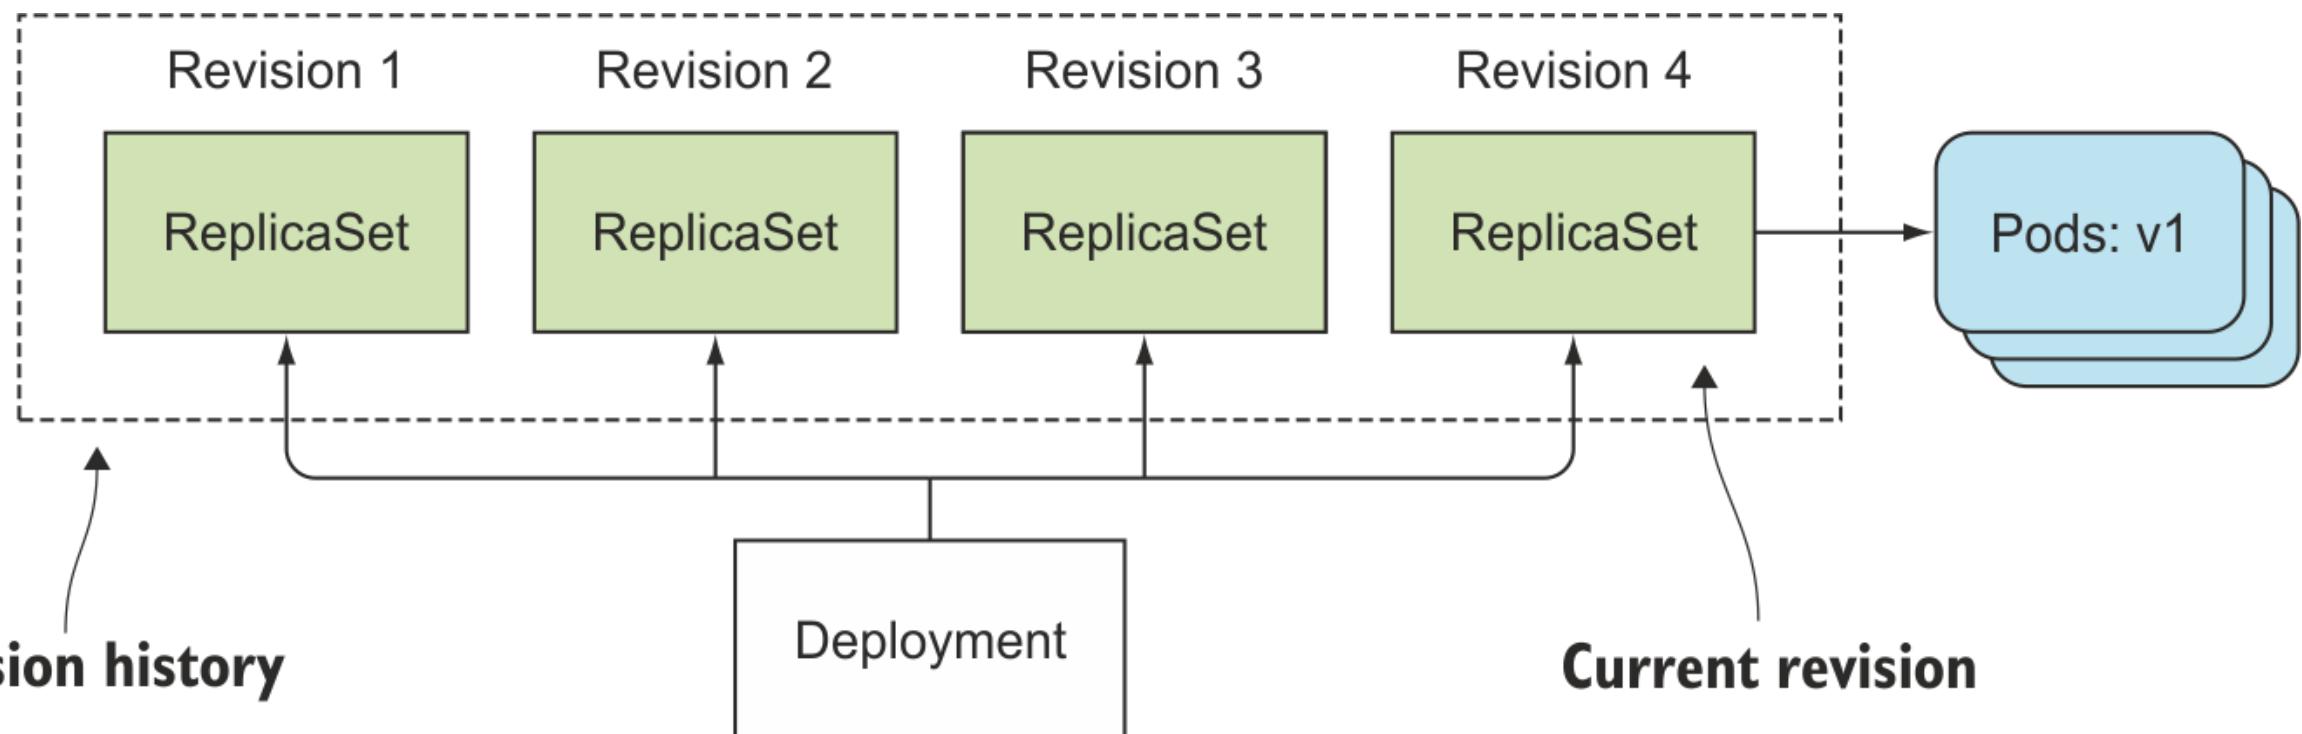
## Rolling back - Simulation

- The revision **history** can be **displayed** with the kubectl rollout history command:  
`$ kubectl rollout history deployment kubia`
- Without --record, the **CHANGE-CAUSE** column in the revision history would be **empty**.
- To roll **back** to a specific **revision** by specifying the revision in the undo command:  
`$ kubectl rollout undo deployment kubia --to-revision=1`

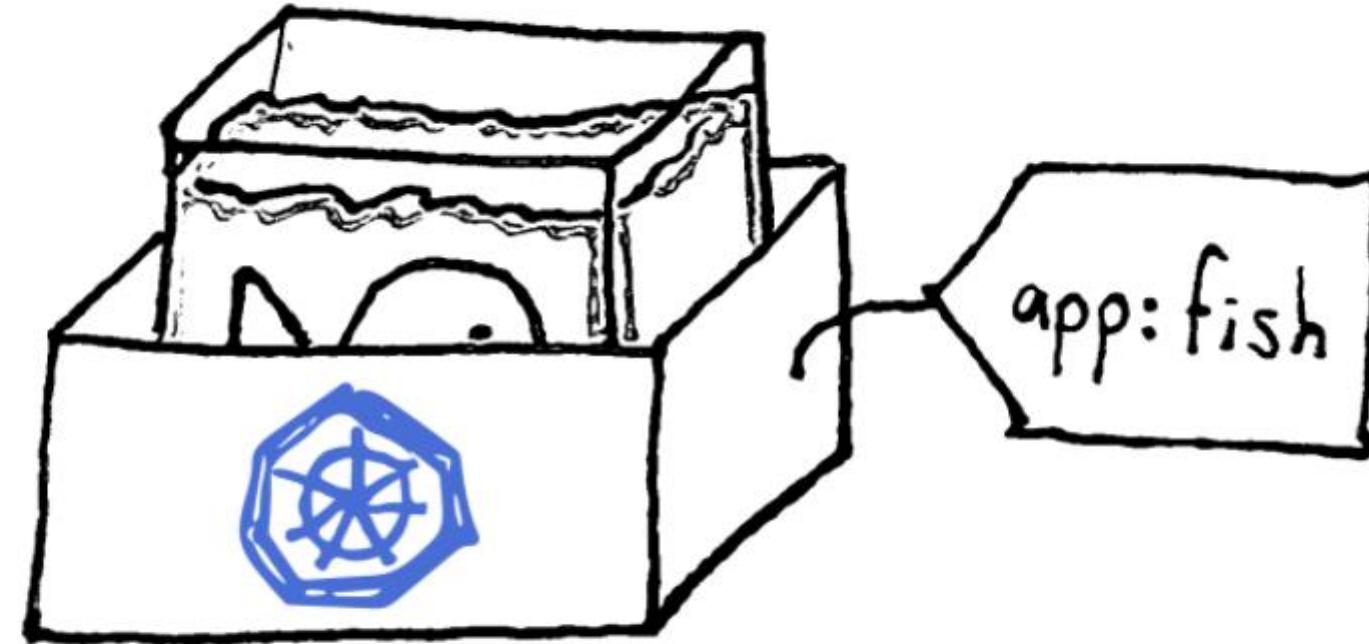
# Deployments



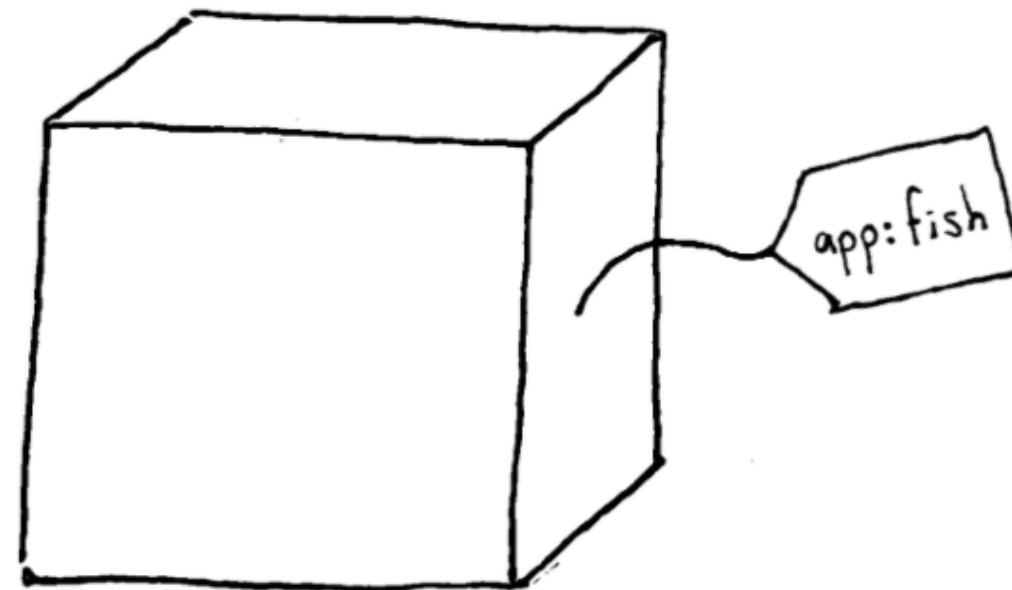
## Rolling back – Simulation

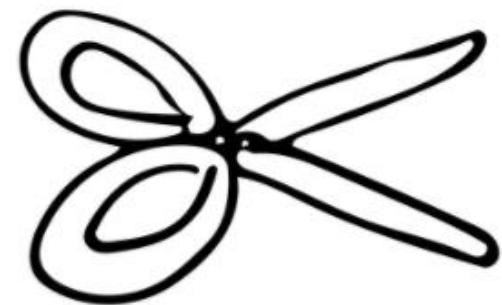
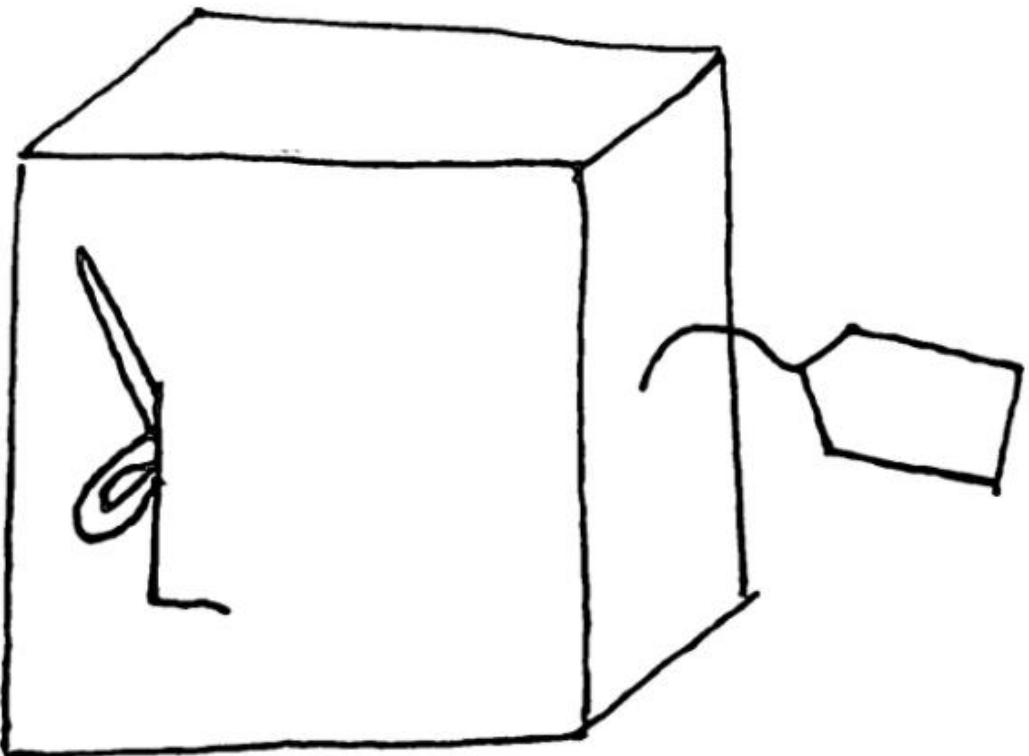


pod



pod





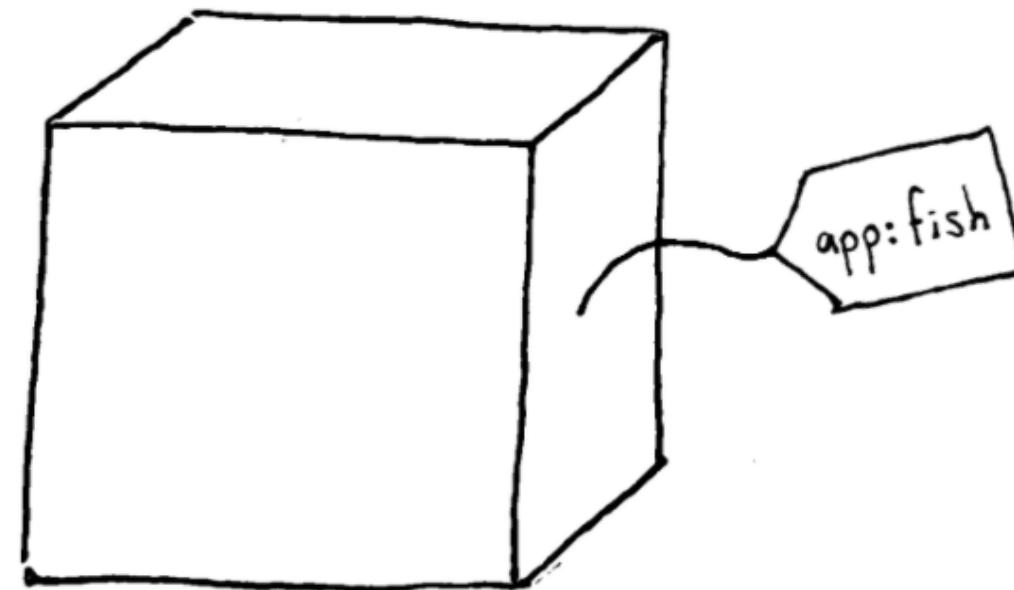
service

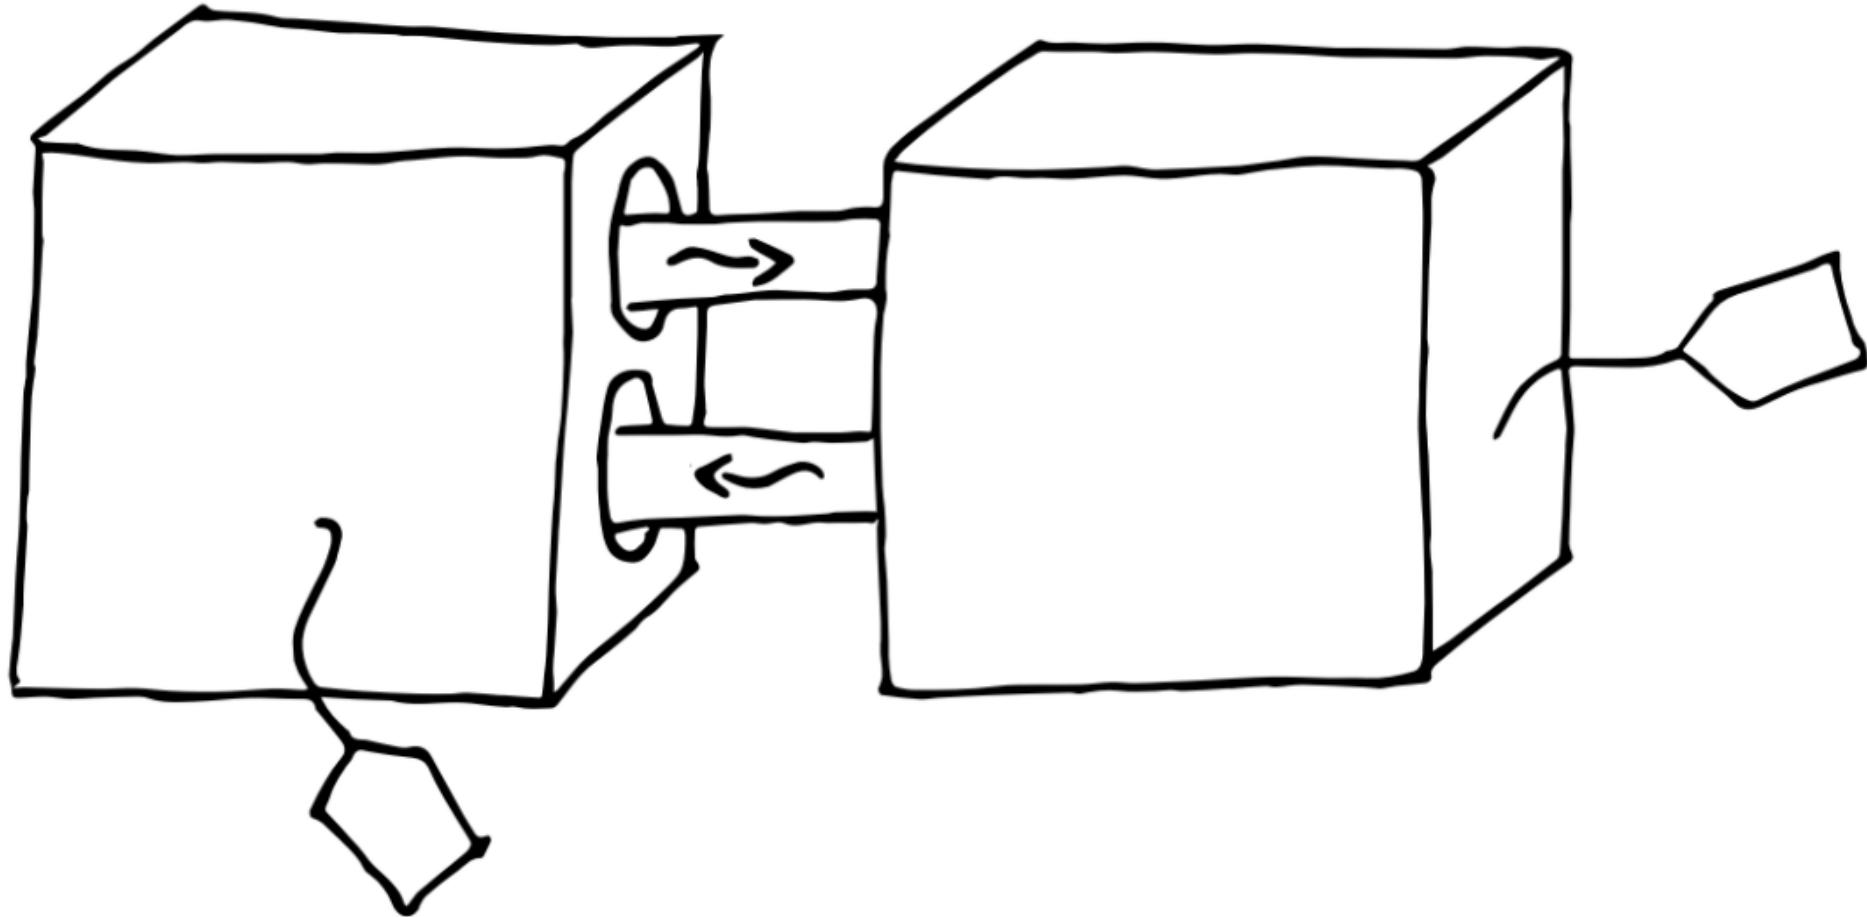
# Services



- A Kubernetes **Service** is a **resource** that **makes a single, constant point of entry to a group of pods** providing the **same service**.
- Each **service has an IP** address and **port** that **never change** while the service exists.
- Clients can **open connections to that IP and port**, and those connections are then **routed to one of the pods** backing that service.

pod





services between pods

# Services



## Service creation – via manifest

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: kubia
```

# Services



## Service creation – via manifest

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: kubia
```

- The port this service will be available on
- The container port the service will forward to

# Services



## Service creation – via manifest

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: kubia
```

- All pods with the `app=kubia` label will be part of this service.

# Services



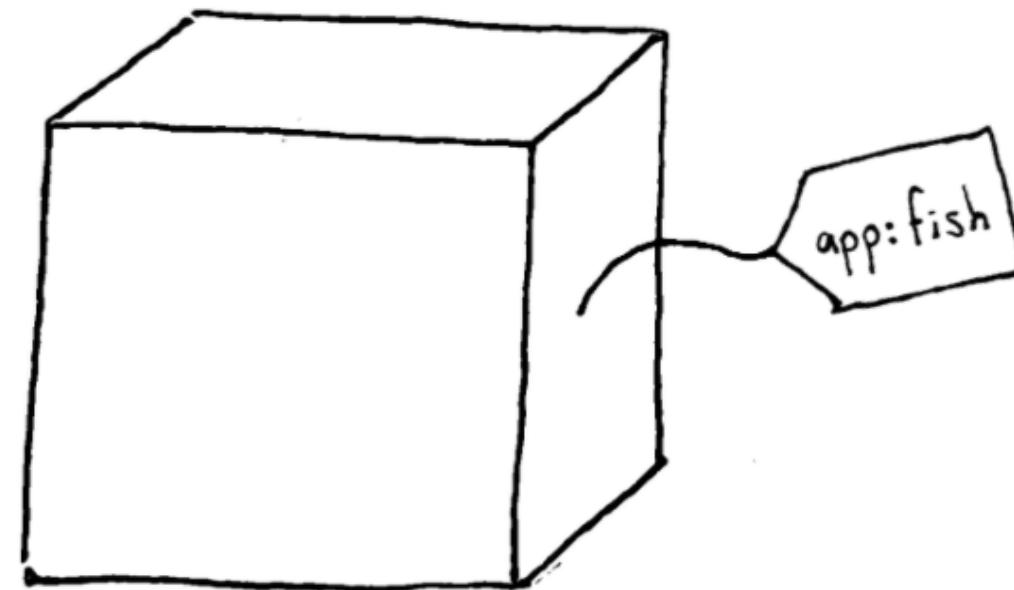
## Service creation – via manifest

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: kubia
```

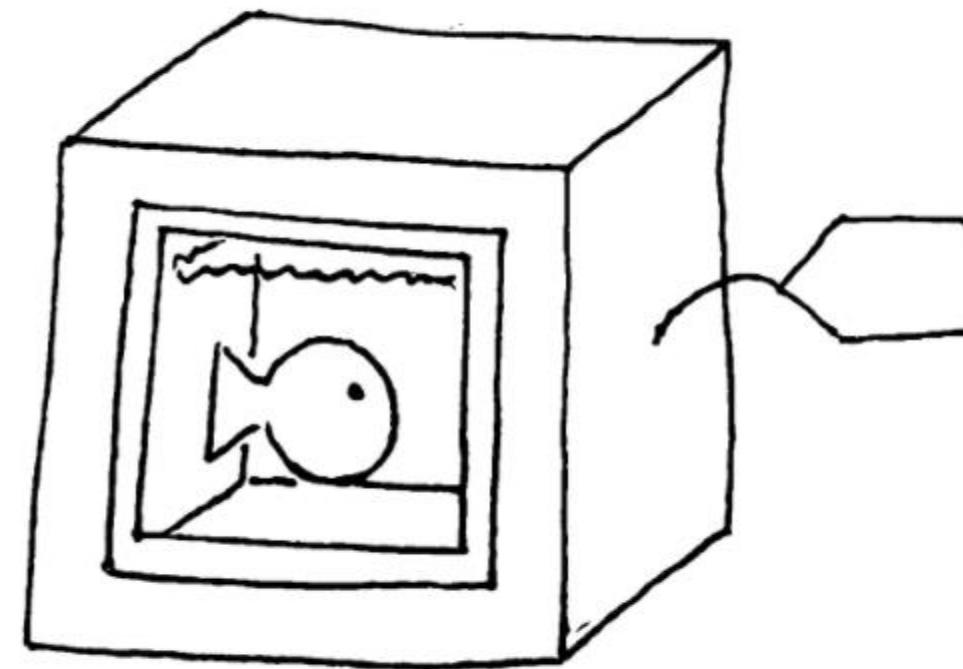
```
$ kubectl create -f kubia-svc.yaml
$ kubectl get svc
```

- This is a **clusterIP** service,
- It's only accessible from inside the cluster.
- The primary purpose of this service is exposing groups of pods to other pods in the cluster

pod



exposed pod



# Services



## NodePort service

- With NodePort, Kubernetes **reserve a port on all its nodes**
  - (the **same port** number is used across all of them)
- Kubernetes will **forward incoming connections** to the **pods** that are **part of the service**.

# Services



## NodePort service

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-nodeport
spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123
  selector:
    app: kubia
```

# Services



## NodePort service

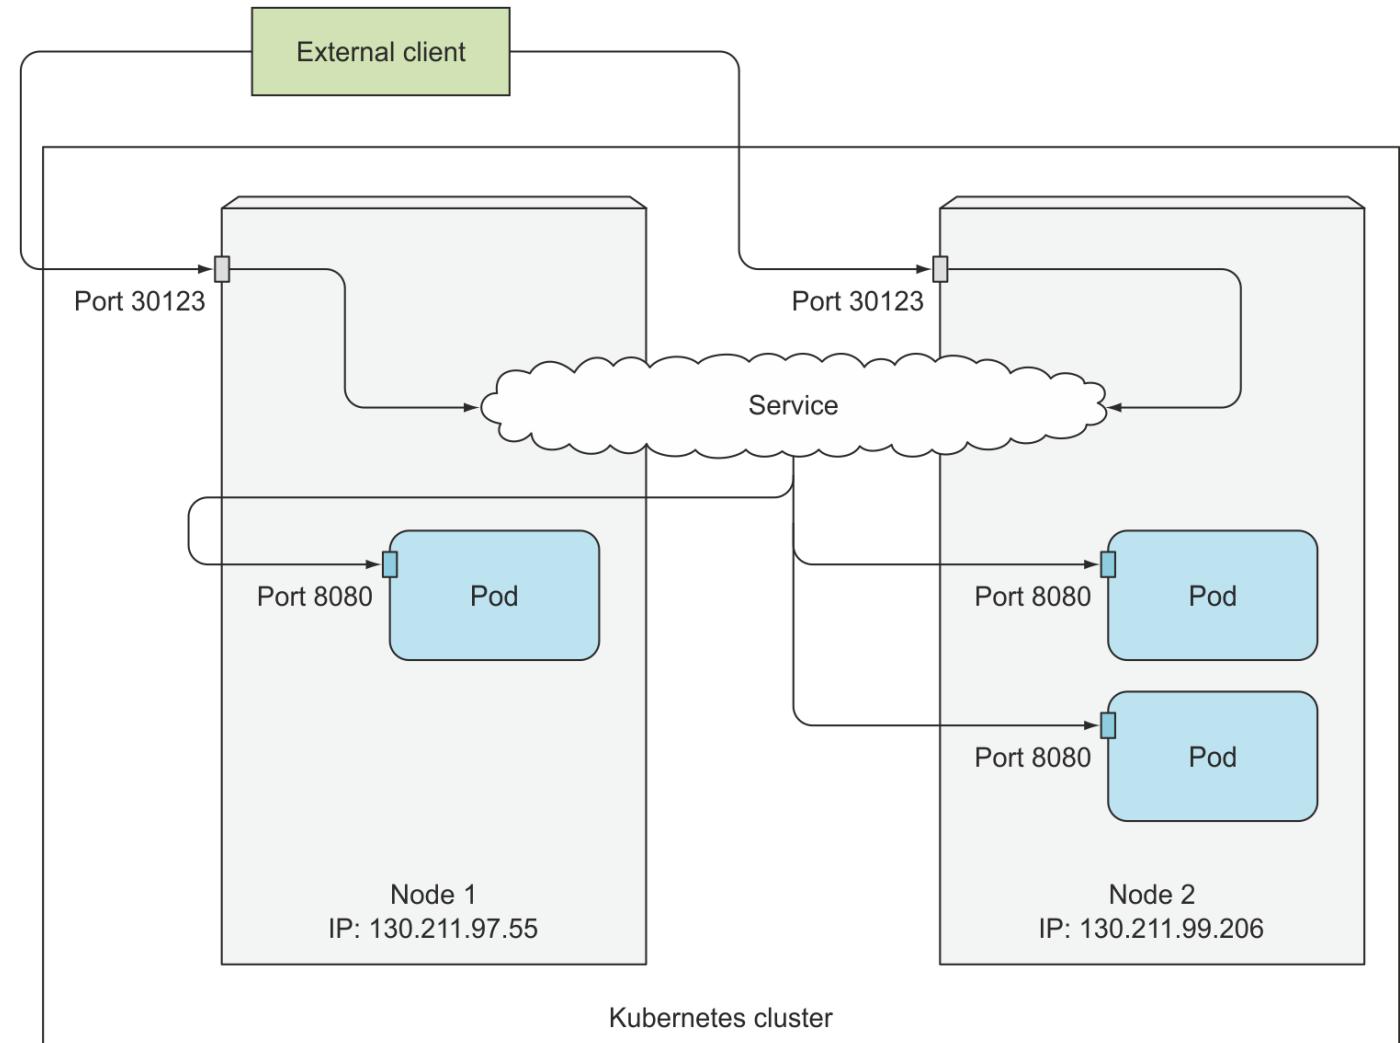
```
apiVersion: v1
kind: Service
metadata:
  name: kubia-nodeport
spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123
  selector:
    app: kubia
```

# Services



## NodePort service

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-nodeport
spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123
  selector:
    app: kubia
```



# Services



## NodePort service

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-nodeport
spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123
  selector:
    app: kubia
```

```
$ kubectl create -f kubia-svc-nodeport.yaml
```

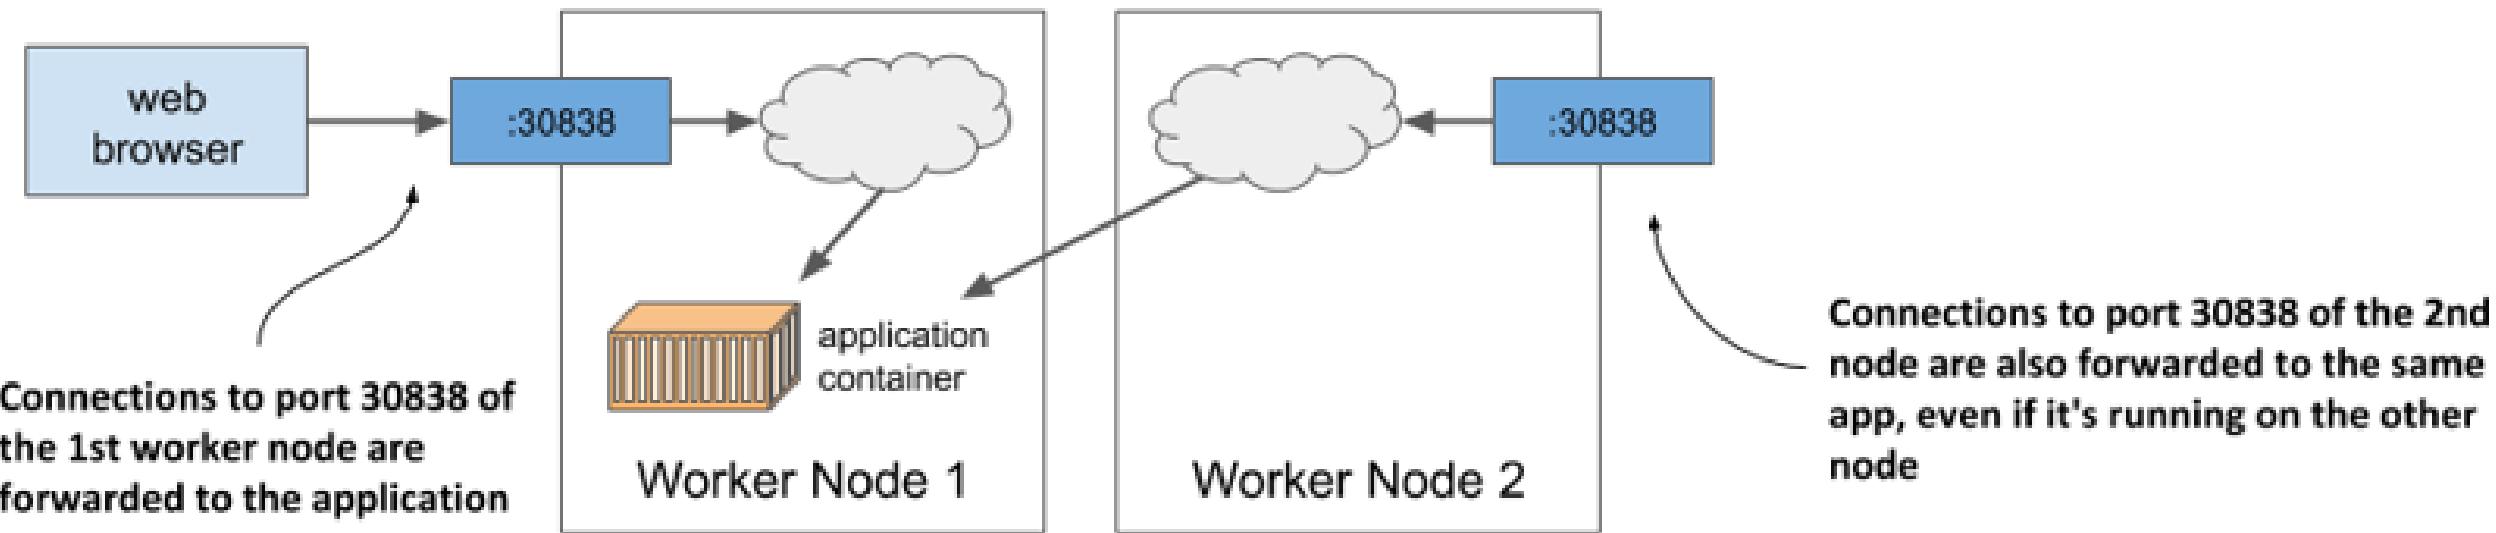
```
$ kubectl get svc kubia-nodeport
```

- The service is **accessible** at the following addresses:
  - <Service IP>: :80
  - <1st node's IP>:30123
  - <2nd node's IP>:30123

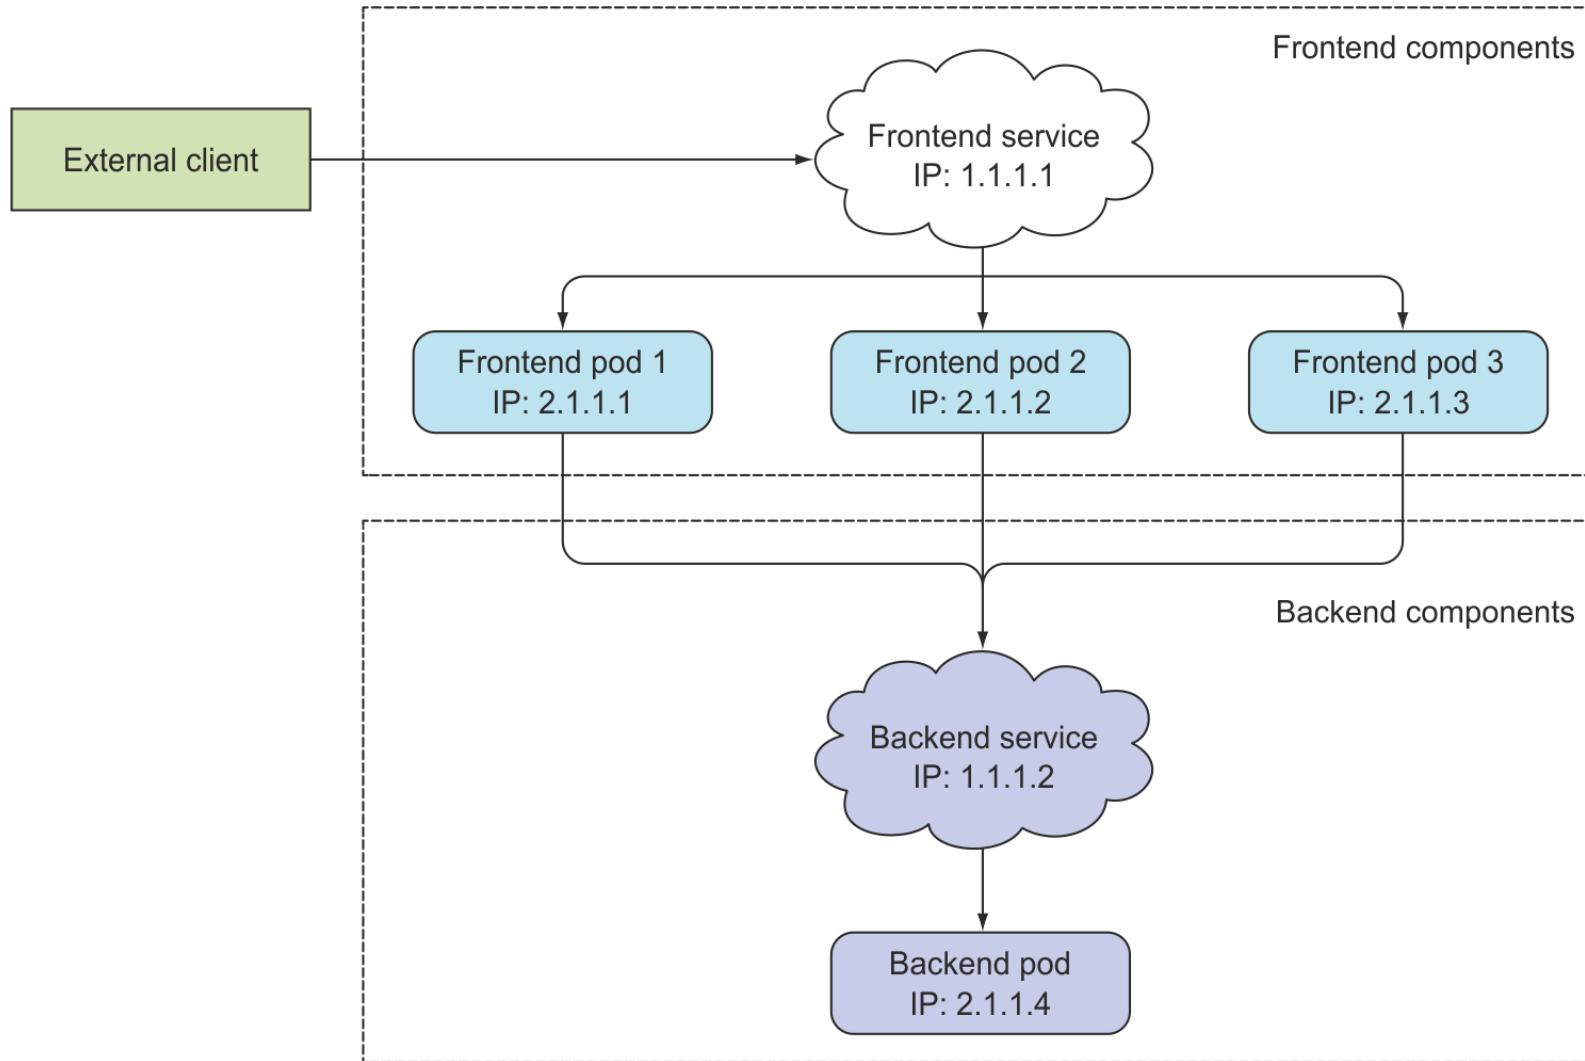
# Services



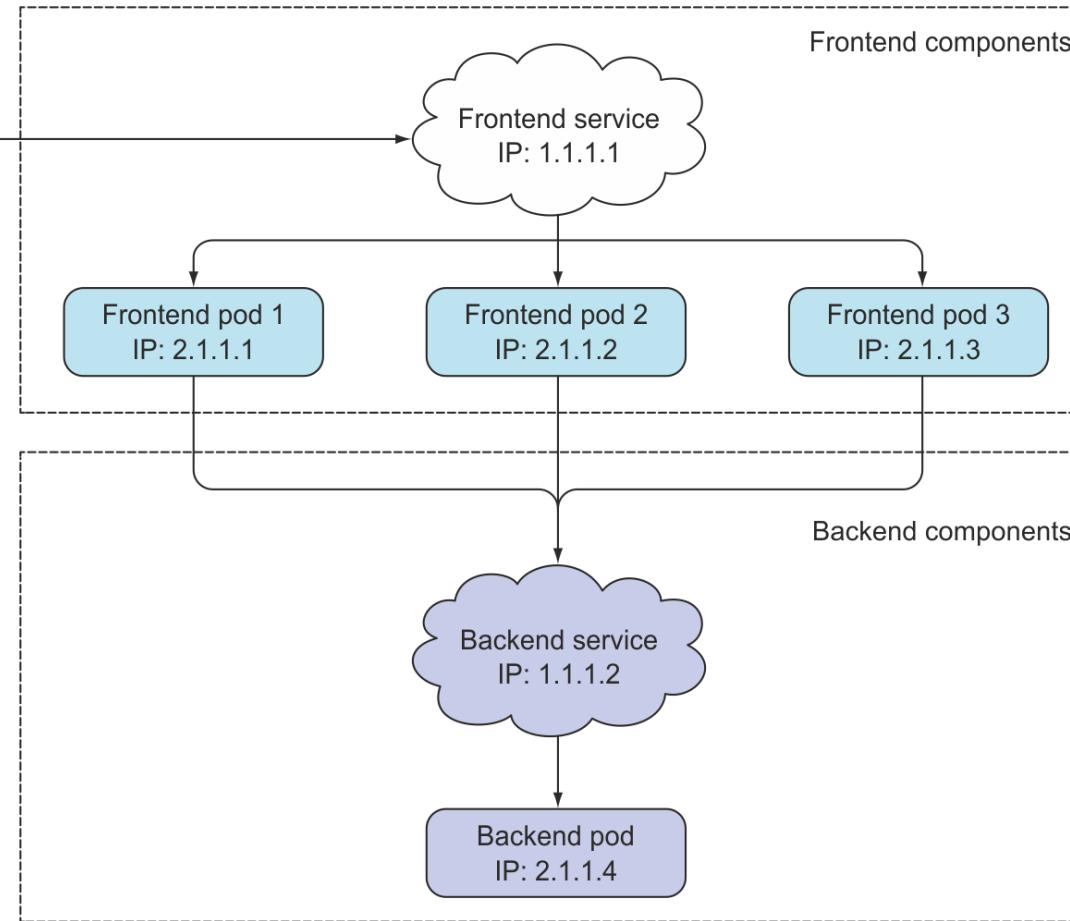
## NodePort service



# Services



# Services



- A **service** created for the **frontend** pods and configuring it to be **accessible from outside** the cluster
  - This will **expose a single, constant IP** address through which **external clients** can **connect** to the pods.
- A **service** created for the **backend** pod,
  - This will create a **stable address** for the **backend** pod.
  - The service **address doesn't change** even if the **pod's IP address changes**

namespace: EDB





## Before Kubernetes and Virtualization

Application

Operating System

Server

## After Kubernetes

Load Balancer

Ingress

Kube-Proxy

Service Mesh

Side Car

Application

Operating System

Virtualization

Server





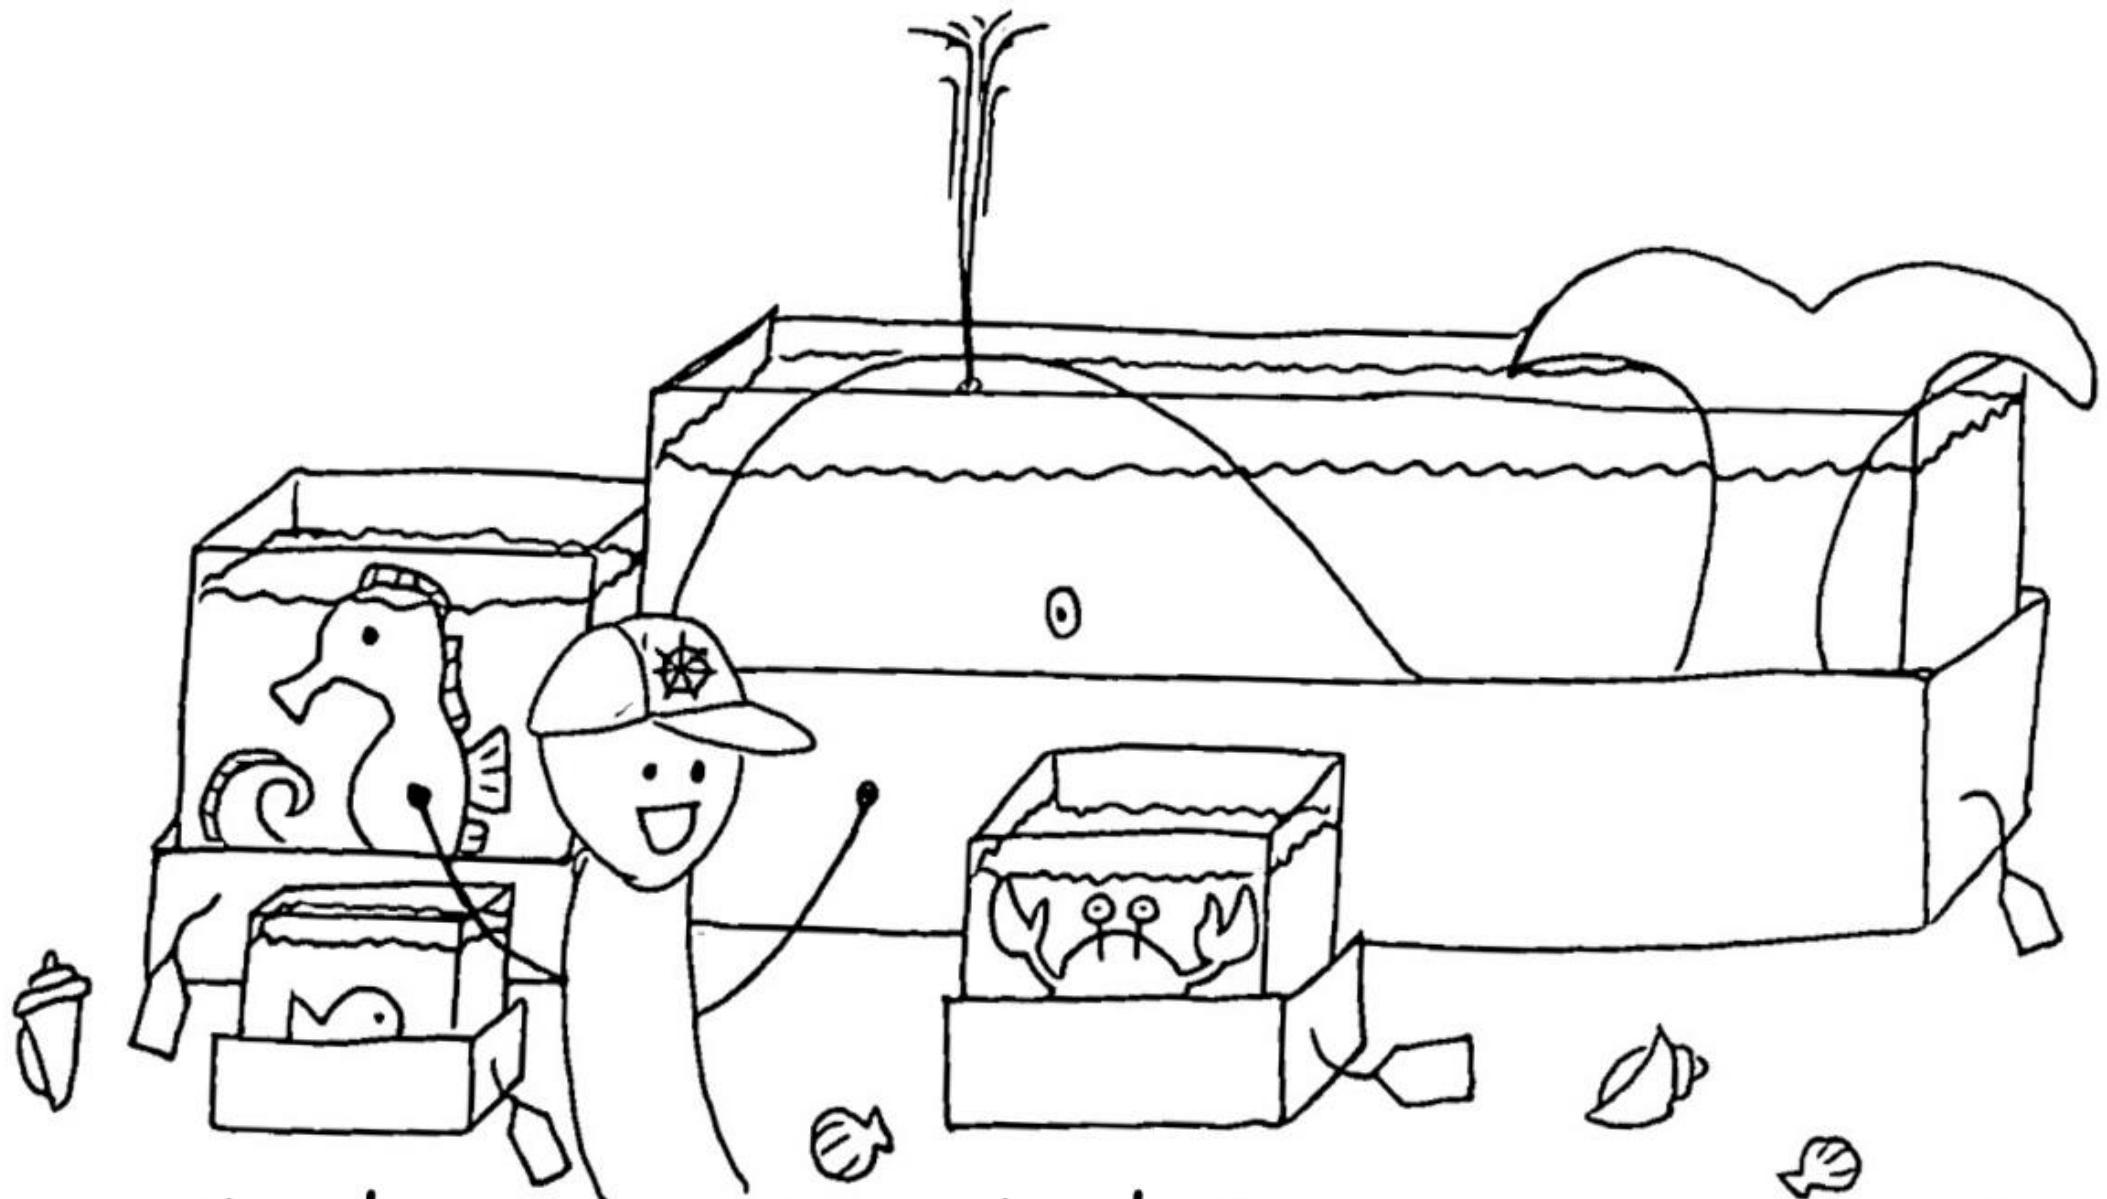
**Running a bare metal  
Kubernetes cluster in  
production isn't  
stressing me out  
anymore.**

**— Mark, 22 years old**





Thank you



thanks for visiting the kubernetes aquarium!

